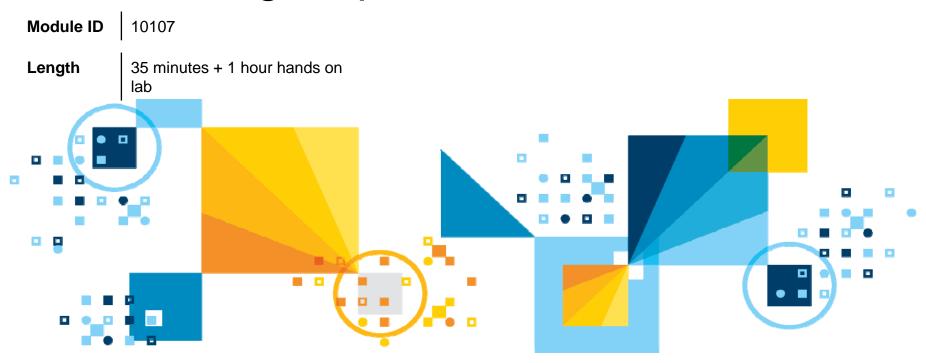


DB2 - Storage Optimizations



For questions about this presentation contact askdata@ca.ibm.com



Disclaimer

© Copyright IBM Corporation 2015. All rights reserved.

THE INFORMATION CONTAINED IN THIS PRESENTATION IS PROVIDED FOR INFORMATIONAL PURPOSES ONLY. WHILE EFFORTS WERE MADE TO VERIFY THE COMPLETENESS AND ACCURACY OF THE INFORMATION CONTAINED IN THIS PRESENTATION, IT IS PROVIDED "AS IS" WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED. IN ADDITION, THIS INFORMATION IS BASED ON IBM'S CURRENT PRODUCT PLANS AND STRATEGY, WHICH ARE SUBJECT TO CHANGE BY IBM WITHOUT NOTICE. IBM SHALL NOT BE RESPONSIBLE FOR ANY DAMAGES ARISING OUT OF THE USE OF, OR OTHERWISE RELATED TO, THIS PRESENTATION OR ANY OTHER DOCUMENTATION. NOTHING CONTAINED IN THIS PRESENTATION IS INTENDED TO, NOR SHALL HAVE THE EFFECT OF, CREATING ANY WARRANTIES OR REPRESENTATIONS FROM IBM (OR ITS SUPPLIERS OR LICENSORS), OR ALTERING THE TERMS AND CONDITIONS OF ANY AGREEMENT OR LICENSE GOVERNING THE USE OF IBM PRODUCTS AND/OR SOFTWARE.

IBM, the IBM logo, ibm.com are trademarks or registered trademarks of International Business Machines Corporation in the United States, other countries, or both. If these and other IBM trademarked terms are marked on their first occurrence in this information with a trademark symbol (® or ™), these symbols indicate U.S. registered or common law trademarks owned by IBM at the time this information was published. Such trademarks may also be registered or common law trademarks in other countries. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at www.ibm.com/legal/copytrade.shtml

Other company, product, or service names may be trademarks or service marks of others.



Module Information

- You should have completed or acquired the necessary knowledge for the following modules in order to complete this module:
 - DB2 Fundamentals
 - DB2 Storage Design
- After completing this module, you should be able to:
 - Describe the storage optimizations available in DB2
 - Explain the benefits of compression
 - Perform the following tasks:
 - Compress database tables
 - Determine the amount of compression is possible for a table



Topics

- Compression Overview
- Storage Optimization in DB2
 - Value Compression
 - Row Compression
 - Compression in BLU
 - Temporary Table Compression
 - XML Compression
 - Index Compression
 - Backup Compression
- Compression Benefits
- Client Success



Compression Overview

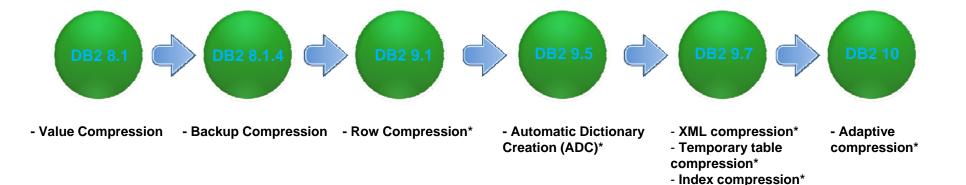
- Requires less storage space and memory
 - Usually the most expensive component in a database solution
- Helps manage storage growth
 - Saves on power and cooling
- Helps improve database performance
 - Fewer I/O operations required
- Reduce administration costs and increase productivity of the IT team
- Facilitates lower cost of ownership





Storage Optimization in DB2

- Provides storage compression services to optimize the performance and footprint of your data
- Basic compression features included in several editions
- Advanced compression features including Adaptive Compression and classic row compression are included in the following versions of DB2:
 - DB2 Advanced Workgroup Sever Edition
 - DB2 Advanced Enterprise Server Edition
 - DB2 Developer Edition



6 © 2015 IBM Corporation

- LOB inlining

^{*} Included in DB2 Advanced Editions (and Developer Edition) as of DB2 10.5



Value Compression

- Removes duplicate entries for a value, only storing one copy
- Can offer savings if your table:
 - has many rows with columns that contain the same value, such as a city name
 - has columns that contain the default value for the data type of the column
- Can be used in combination of either adaptive row compression or classic row compression
- Enabled with VALUE COMPRESSION option in CREATE TABLE
- Compress system default column values with clause COMPRESS SYSTEM DEFAULT

```
CREATE TABLE EMPLOYEE_SALARY

(DEPTNO CHAR(3) NOT NULL,

EMPNO CHAR(6) NOT NULL,

SALARY DECIMAL(9,2) NOT NULL WITH DEFAULT COMPRESS SYSTEM DEFAULT)

VALUE COMPRESSION;
```



Row Compression

- Also known as deep compression
- Uses a dictionary-based compression algorithm to replace recurring strings with shorter symbols within rows

compression*

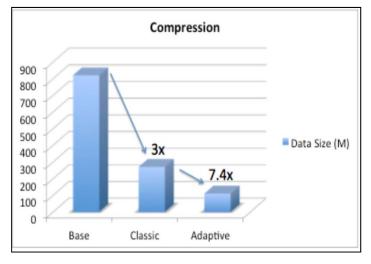
- Index compression*

Continuous enhancement since it was introduced in DB2 9.1



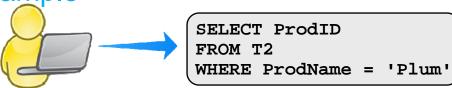
- Two types available:
 - LOB inlining - Classic (static) row compression
 - Adaptive row compression
 - An enhancement to classic row compression to provide extra storage savings

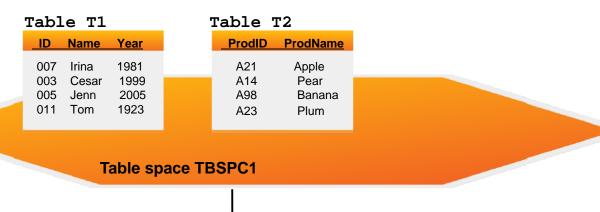


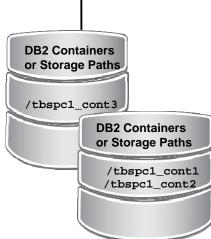




Understanding the Basics The DB2 Storage - Example



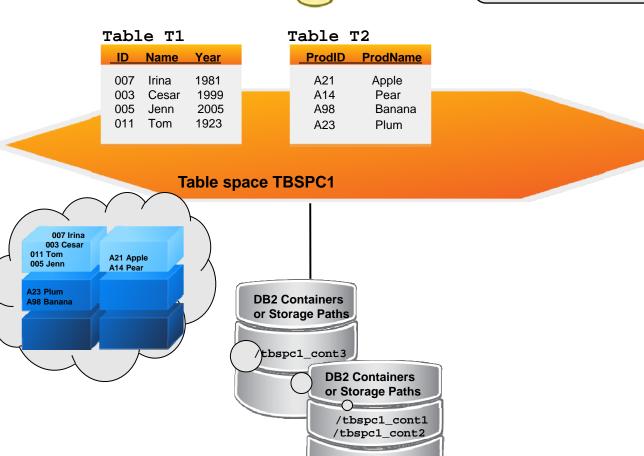






Understanding the Basics The DB2 Storage - Example





11



© 2015 IBM Corporation

Understanding the Basics The DB2 Storage - Example SELECT ProdID FROM T2 WHERE ProdName = 'Plum' Table T1 Table T2 **DB2 Buffer Pool** Year **ProdID ProdName** Name A21 Apple A14 Pear A98 Banana 007 Irina 1981 A21 Apple A23 Plum 003 Cesar 1999 A14 Pear 005 Jenn 2005 A98 Banana 011 Tom 1923 A23 Plum **Table space TBSPC1** 007 Irina 003 Cesar 011 Tom A21 Apple 005 Jenn A14 Pear A23 Plum Disk **DB2 Containers** or Storage Paths **Extent Data** tbspc1_cont3 **Page DB2 Containers Table space** or Storage Paths /tbspc1_cont1 **Table** /tbspc1_cont2

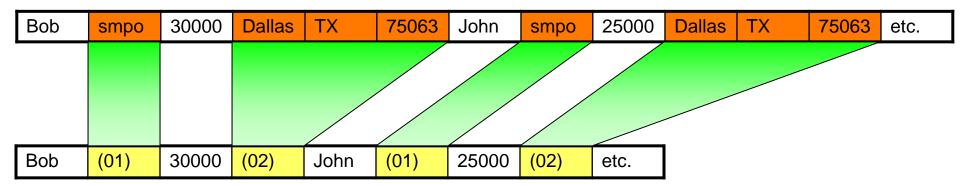


Row Compression - Classic

- Also referred to as static row compression
- Uses a table-level compression dictionary (1 dictionary per table) to compress data by row, across multiple columns
- Dictionary is used to map repeated byte patterns to smaller symbols. These smaller symbols replace long patterns in table rows.
- After dictionary is created, data is compressed as it is inserted/updated in the table.
 - DB2 automatically creates the dictionary when enough the table has enough data for sampling

Name	Dept	Salary	City	ST	ZIP
Bob	smpo	30000	Dallas	TX	75063
John	smpo	25000	Dallas	TX	75063

Dictionary				
(01)	smpo			
(02)	Dallas, TX, 75063			





Why Choose Adaptive Row Compression

- Classic row compression works very well
 - Very fast and robust
 - Not sensitive to data clustering and ordering
- But it has limitations
 - Dictionary requires classic table reorganization to refresh
 - Not sensitive to data clustering and ordering
 - Dictionary capacity limits compression ratios for some large tables
- DB2 10 introduces Adaptive Compression
 - An enhancement to classic row compression
 - Industry-leading compression technology for row stores
- Reduces TCO, providing extra compression savings:
 - Classic row compression: Typically saves ~40%-75%
 - Adaptive row compression: Typically saves ~75%-85%
 - Adaptive typically saves 30% over classic row compression

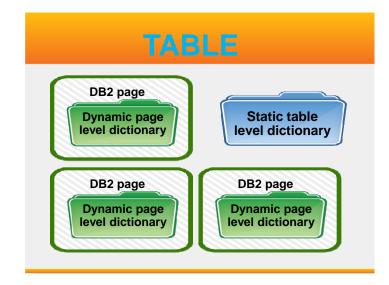


Row Compression – Adaptive

- Adaptive Compression is an enhancement to the Classic Row Compression
 - Compress rows by using a combination of two types of algorithms
 - Globally recurring byte sequences
 - Locally recurring byte sequences

Benefits

- Page level dictionaries adapt to data skew over a period of time
- No REORGs required to maintain high compression ratio
- Less disk space for data and logs
- Reduced I/O. Fewer pages to process



- Better compression ratios than Classic Row Compression
- Over time reduces need of a REORG as page-level dictionaries adapt to data skew over time (i.e. maintain compression rates)
- Default row compression method in DB2 10





[1] California 9

How Does Adaptive Compression Work?

Step 1: Compression with static table level dictionary

Christine	Haas	(408) 463-1234	555 Bailey Avenue	San Jose	California	95141
John	Thompson	(408) 463-5678	555 Bailey Avenue	San Jose	California	95141
Jose	Fernandez	(408) 463-1357	555 Bailey Avenue	San Jose	California	95141
Margaret	Miller	(408) 463-2468	555 Bailey Avenue	San Jose	California	9 5141
Bruce	Kwan	(408) 956-9876	4400 North 1st Street	San Jose	California	9 5134
James	Geyer	(408) 956-5432	4400 North 1st Street	San <mark>Jose</mark>	California	9 5134
Linda	Hernandez	(408) 956-9753	4400 North 1st Street	San <mark>Jose</mark>	California	9 5134
Theodore	Mills	(408) 927-8642	650 Harry Road	San Jose	California	9 5134
Susan	Stern	(408) 927-9630	650 Harry Road	San Jose	California	9 5134
James	Polaski	(415) 545-1423	425 Market Street	San Francisco	California	94105
John	Miller	(415) 545-5867	425 Market Street	San Francisco	California	94105
James	Walker	(415) 545-4132	425 Market Street	San Francisco	California	94105
Elizabeth	Brown	(415) 545-8576	425 Market Street	San Francisco	California	94105
Sarah	Johnson	(415) 545-1928	425 Market Street	San Francisco	California	94105

[2] San
[3] Jose
[4] Francisco
[5] Avenue
[6] Street
[7] Road

Compression with global table static dictionary

- Table level compression symbol dictionary containing globally recurring patterns
- Table-level dictionary can only be rebuilt during Classic Table REORG
 - Involves re-compressing all data

Christine	Haas	(408) 463-1234	555 Bailey [5]	[2][3]	[1]	95141
John	Thompson	(408) 463-5678	555 Bailey [5]	[2][3]	[1]	95141
[3]	Fernandez	(408) 463-1357	555 Bailey [5]	[2][3]	[1]	95141
Margaret	Schneider	(408) 463-2468	555 Bailey [5]	[2][3]	[1]	95141
Bruce	Kwan	(408) 956-9876	4400 North 1st [6]	[2][3]	[1]	95134
James	Geyer	(408) 956-5432	4400 North 1st [6]	[2][3]	[1]	95134
Linda	Hernandez	(408) 956-9753	4400 North 1st [6]	[2][3]	[1]	95134
Theodore	Mills	(408) 927-8642	650 Harry [7]	[2][3]	[1]	95134
Susan	Stern	(408) 927-9630	650 Harry [7]	[2][3]	[1]	95134
James	Polaski	(415) 545-1423	425 Market [6]	[2][4]	[1]	94105
John	Miller	(415) 545-5867	425 Market [6]	[2][4]	[1]	94105
James	Walker	(415) 545-4132	425 Market [6]	[2][4]	[1]	94105
Elizabeth	Miller	(415) 545-8576	425 Market [6]	[2][4]	[1]	94105
Sarah	Johnson	(415) 545-1928	425 Market [6]	[2][4]	[1]	94105



How Does Adaptive Compression Work?

Step 2: Compression w/ Page-Level Dictionaries

	Christine	Haas	(408) 463-1234	555 Bailey [5]	[2][3]	[1]	5141
<u>e</u>	John	Thompson	(408) 463-5678	555 Bailey [5]	[2][3]	[1]	5141
page	Ellen	Fernandez	(408) 463-1357	555 Bailey [5]	[2][3]	[1]	5141
	Margaret	Schneider	(408) 463- <mark>2</mark> 468	555 Bailey [5]	[2][3]	[1]	5141
Data	Bruce	Kwan	(408) 956-9876	4400 North 1st [6]	[2][3]	[1]	5134
	James	Geyer	(408) 956-5432	4400 North 1st [6]	[2][3]	[1]	5134
	Linda	Hernandez	(408) 956-9753	4400 North 1st [6]	[2][3]	[1]	5134

1)	Theodore	Mills	(408) 927-8642	650 Harry [7]	[2][3]	[1]	5134
2	Susan	Stern	(408) 927-9630	650 Harry [7]	[2][3]	[1]	5134
2	James	Polaski	(415) 545-1423	425 Market [6]	[2][4]	[1]	4105
7	John	Miller	(415) 545-5867	425 Market [6]	[2][4]	[1]	4105
9	James	Walker	(415) 545-4132	425 Market [6]	[2][4]	[1]	4105
-	Elizabeth	Miller	(415) 545-8576	425 Market [6]	[2][4]	[1]	4105
	Sarah	<mark>John</mark> son	(415) 545-1928	425 Market [6]	[2][4]	[1]	4105



Christ	tine Haas	(2)12	34 (4)
John	Thom	pson (2)56	78 (4)
Ellen	F(1)	(2)13	57 (4)
Marga	aret Schn	eider (2)24	68 (4)
Bruce	. Kwan	(3)98	76 (5)
James	s Geyei	r (3)54	32 (5)
Linda	H(1)	(3)97	53 (5)

(1)	ernandez
(2)	(408) 463-
(3)	(408) 956-
(4)	555 Bailey [5] [2][3] [1] 5141
(5)	4400 North 1st [6] [2][3] [1] 5134
	Б

Page level dictionary

Theodore	(3)s	(4)8642	(6)
Susan	Stern	(4)9630	(6)
(1)	Polaski	(5)1423	(7)
(2)	(3)er	(5)5867	(7)
(1)	Walker	(5)4132	(7)
Elizabeth	(3)er	(5)8576	(7)
Sarah	(2)son	(5)1928	(7)

(1)	James
(2)	John
(3)	Mill
(4)	(408) 927-
(5)	(415) 545-
(6)	650 Harry [7] [2][3] [1] 5134
(7)	425 Market [6] [2][4] [1] 4105

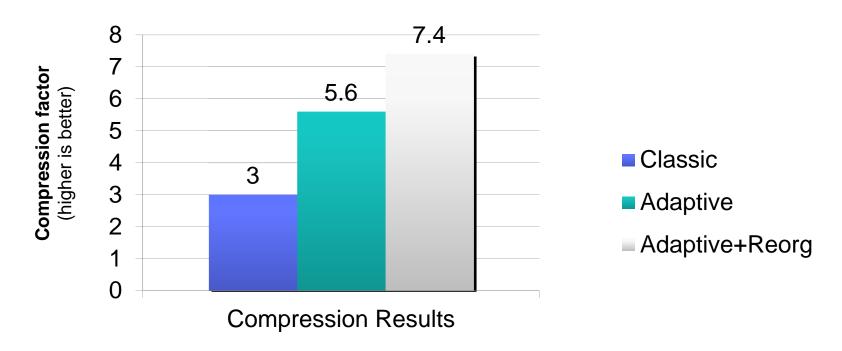
Page level dictionary

- Page-level compression dictionaries contain locally frequent patterns
- Page-level compression dictionary building & rebuilding is fully automatic
- Algorithm optimized for compressing data already compressed by table-level dictionary
- Page-level compression dictionaries are stored as special records in each page



Data Warehouse Compression Results 230GB raw size - Most of the data in a single table

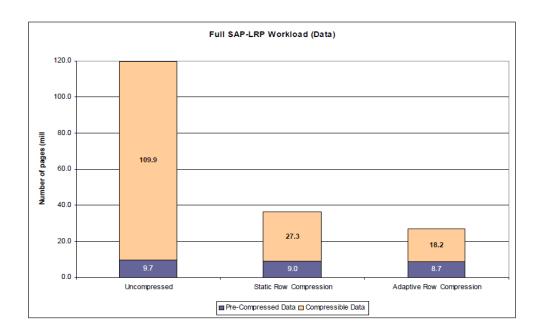
- Graph Storage Savings
- Increase in savings by Adaptive Compression
 - 3x Compression with Static Compression using reorg
 - 5.6x Compression with Automatic dictionary creation and Adaptive Compression
 - 7.4x Compression with Adaptive Compression and full reorg





Additional Storage Savings

SAP-LRP workload, ~2TB data size before compression

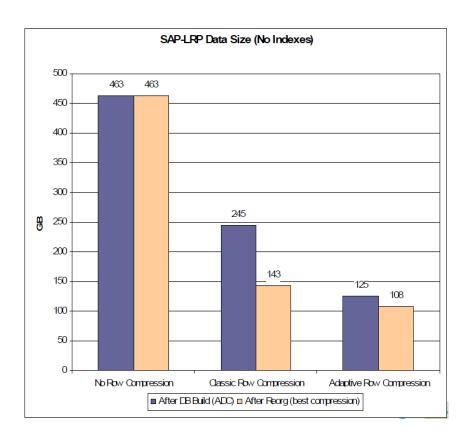


- Considering the 1425 largest tables (~99.5% of the total data), after Classic Table Reorg
- 33% savings for normal tables, compared to Classic Row Compression
- Almost no storage savings for tables with pre-compressed data (SAP cluster tables, certain tables with LOB fields)



Reorg Avoidance – ADC vs. Classic Table Reorg

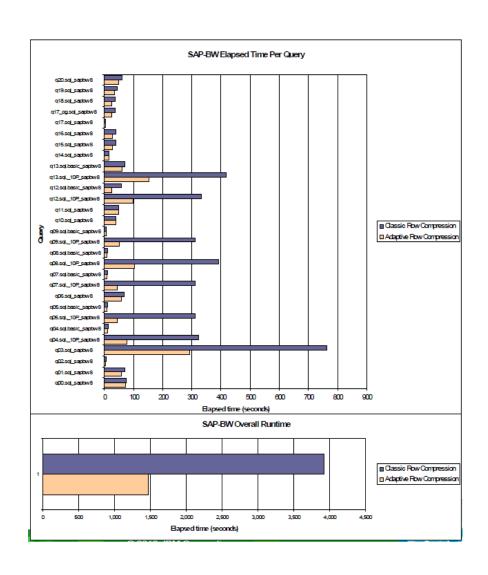
- SAP-LRP system with ~450GB of data in uncompressed form
- ~24,000 tables, 99.7% of the data in 1,500 tables
- Comparison between database build (ADC) and full Classic Table Reorg
- Classic Row Compression with ADC yields
 71% larger data size than after Reorg
- Adaptive Compression with ADC is only 16% larger than after Reorg
- Adaptive Compression with ADC is 14% smaller than Classic Row Compression after Reorg (i.e. best possible compression)





Query Performance – SAP-BW Benchmark (1TB)

- Additional CPU consumption for compression & expansion operations
 - Row compression design paradigm:
 Compression under logical I/O
 - Rule of thumb: Compression is beneficial in I/O-bound situations
- Performance measurements done from two angles:
 - 1. Focus on additional CPU cost (to minimize overhead)
 - 2. Elapsed time studies in benchmark scenarios
- Most long-running queries experience dramatic speed-up
- Overall runtime reduced by 62%
- Median query speedup of 43%





Row Compression – Enablement & Tools

- How to enable row compression?
 - Must have DB2 Storage Optimization Feature
 - To enable classic row compression

```
CREATE TABLE / ALTER TABLE ... COMPRESS YES STATIC
```

- To enable adaptive row compression

```
CREATE TABLE / ALTER TABLE ... COMPRESS YES
```

To disable compression

```
CREATE TABLE / ALTER TABLE ... COMPRESS NO
```

Adaptive is the default in DB2 10

- Data is compressed after the table dictionary is created.
 - INSERT/UPDATE/LOAD/IMPORT can trigger the automatic dictionary creation
 - Classic REORG with RESETDICTIONARY option will always generate a new dictionary and compress all table data



Row Compression - Example Scenarios

1) Compressing data for new table

CREATE TABLE Sales (<columns definition>) COMPRESS YES

Automatic Dictionary Creation (ADC) will kick off and create compression dictionary with the load command. Once the dictionary is built, new data put into the table is compressed:

LOAD FROM file OF DEL REPLACE INTO NewSale

2) Compressing data in existing tables

ALTER TABLE Sales COMPRESS YES

Existing data is still un-compressed. Explicitly compress existing data via REORG:

REORG TABLE Sales

3) Recreating the dictionary to optimize compression (Classic Row Compression) Data has changed a lot so current dictionary is not so effective anymore. Use REORG to recreate dictionary and re-compress data:

REORG TABLE Sales RESETDICTIONARY

4) Uncompressing your data

Disable compression:

ALTER TABLE Sales COMPRESS NO

Uncompress data:

REORG TABLE Sales

Adaptive Compression greatly reduces the need for REORGs to maintain the compression ratio.





Row Compression – Enablement & Tools

- Estimating storage savings
 - ADMIN_GET_TAB_COMPRESS_INFO
 - ADMIN GET TAB DICTIONARY INFO
 - ADMIN_GET_TAB_COMPRESS_INFO_V97 ← ADMIN_GET_TAB_COMPRESS_INFO_V97 ←

```
SELECT SUBSTR(TABNAME, 1, 10) tabname, OBJECT_TYPE, ROWCOMPMODE,
PCTPAGESSAVED_CURRENT current, PCTPAGESSAVED_STATIC with_static,
PCTPAGESSAVED_ADAPTIVE with_adaptive
FROM TABLE(SYSPROC.ADMIN GET TAB COMPRESS INFO('DB2INST1','CUSTOMERS')) AS T;
```

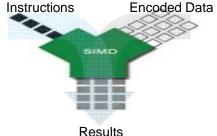


TABNAME	OBJECT_TYPE	ROWCOMPMODE	CURRENT	WITH_STATIC	WITH_ADAPTIVE
CUSTOMERS	DATA	s	60	68	81
CUSTOMERS	XML	S	58	62	62



BLU uses Multiple Compression Techniques

- Approximate Huffman-Encoding ("frequency-based compression"), prefix compression, and offset compression
- For column-organized tables, COMPRESSION cannot be enabled/disabled
- Frequency-based compression: Most common values use fewest bits
 - Exploiting skew in data distribution improves compression ratio
 - Very effective since all values in a column have the same data type
 - Maps entire values to dictionary codes
- Actionable compression: Order-preserving encoding allows predicates to be evaluated on compressed data
- SIMD (Single Instruction Multiple Data) parallelism used for fast predicate evaluation on multiple compressed values
- Avoiding decompression during predicate evaluation provides significant query performance gains





Temporary Table Compression

- Enabled automatically with the DB2 Storage Optimization Feature
- Only classic row compression is used for temporary tables
- System temporary tables
 - DB2 Optimizer analyzes storage savings and performance impact, if it is worthwhile to compress, classic row compression is used automatically
- User temporary tables
 - Created global temporary tables (CGTTs) and declared global temporary tables (DGTTs) are always enabled for compression
- Use db2pd to see whether the optimizer used compression for system temporary tables

db2pd -db <database> -temptable



Index Compression

- Especially useful for large OLTP and data warehouse environments
- Index is compressed by default for compressed tables
- Can specify COMPRESS YES option in CREATE INDEX
- Different algorithms implemented by the DB2 engine for:
 - RID list compression
 - Prefix compression
 - Variable slot directory

CREATE INDEX <index name> COMPRESS YES

ALTER INDEX <index_name> COMPRESS YES
REORG INDEXES ALL FOR TABLE <table_name>



Index Compression – Monitoring

■ How much space could I save by compressing the indexes on table T1?

```
SELECT index_name, pages_saved_percent, compress_attr, index_compressed FROM TABLE(SYSPROC.ADMIN_GET_INDEX_COMPRESS_INFO ('T', 'myschema', 'T1', NULL, NULL)) AS T
```



INDEX_NAME		ENT_PAGES_SAVED	COMPRESS_ATTR	IN.
INDEX1	57	N	N	

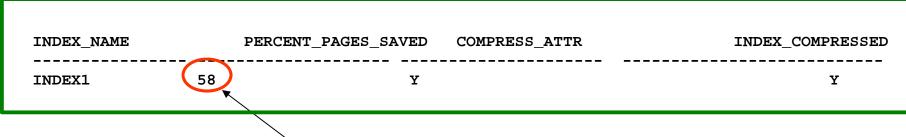


Index Compression – Monitoring

■ How much space did I save by compressing the indexes on table T1?

```
ALTER INDEX index1 COMPRESS YES
REORG INDEXES ON TABLE t1
RUNSTATS ON TABLE t1
SELECT index_name, pages_saved_percent, compress_attr, index_compressed
FROM TABLE (SYSPROC.ADMIN_GET_INDEX_COMPRESS_INFO('T', 'myschema', 'T1', NULL, NULL)) AS T
```



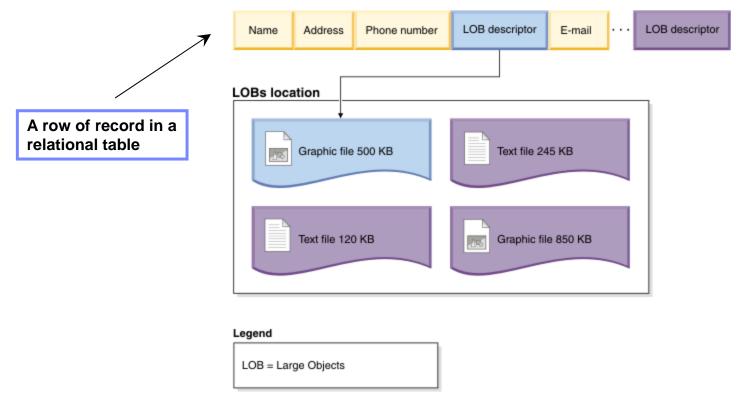


actual savings



XML and LOB Inlining

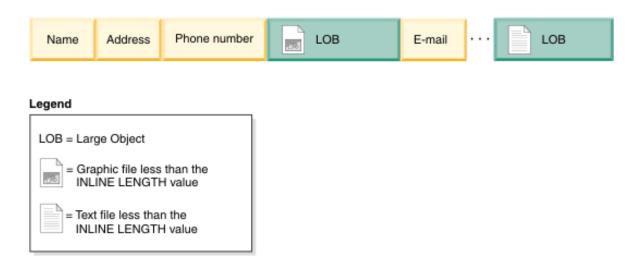
- Default storage behavior
 - All XML documents stored in XML storage object, also referred to as XDA (XML Data Area)
 - All LOB (Large Objects) stored in LOB storage object
 - Descriptors in rows identify location of each object





XML and LOB Inlining

- Inlining
 - Allows small (<32KB for a 32KB page size) XML documents and LOBs to be stored in the base table row
 - Better performance and reduced storage needs



CREATE TABLE PROJECTS (PID INTEGER, PLAN XML INLINE LENGTH 300, STARTDATE DATE, ...)

ALTER TABLE PROJECTS ALTER COLUMN PLAN SET INLINE LENGTH 1004



XML and LOB Compression

- XML (XDA) Compression
 - Available since DB2 9.7
 - Same dictionary approach used for Table Compression
 - Enablement is via the table COMPRESS YES option
 - Classic/'Offline' REORG table based
 - Use LONGLOBDATA option to compress XML data

CREATE TABLE mytabl COMPRESS YES
ALTER TABLE mytabl COMPRESS YES

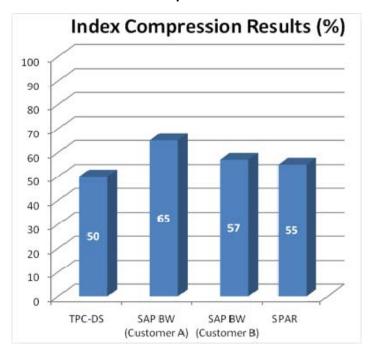
- Inlining
 - Inlined XML or LOB data can be compressed with both adaptive and classic row compression

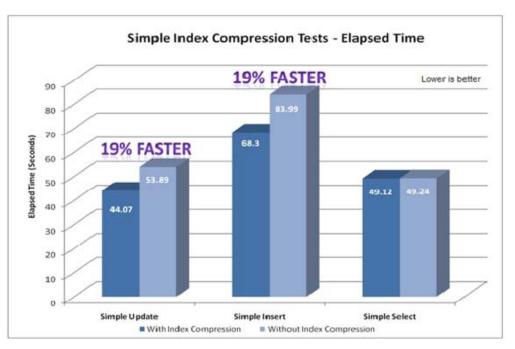




Compression Benefits Index Compression

- Fewer index levels
 - Fewer logical and physical I/Os for key search (insert, delete, select)
 - Better buffer pool hit ratio
- Fewer index leaf pages
 - Fewer logical and physical I/Os for index scans
 - Fewer splits
 - Better buffer pool hit ratio



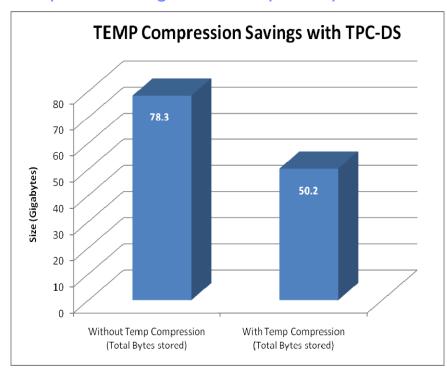




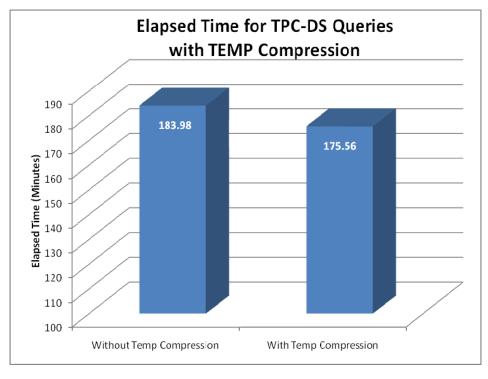
Compression Benefits Temp Table Compression

- Only vendor in industry to compress temp tables
 Internal (DB2 utilities, SORT) and external (DGTTs) compressed
- Savings and performance benefits require no DBA intervention

Space Savings with Temp Compression



Time Savings with Temp Compression





Backup Compression

- Reduce size of database backups in addition to other table-level compressions
- All of data in backup image is compressed, including
 - Catalog tables
 - Index objects
 - LOB objects
 - Auxiliary database files
 - Database meta-data
 - more
- Enable with COMPRESS option in BACKUP DATABASE

BACKUP DATABASE <database> ... COMPRESS



Archive Log Compression

- New in DB2 10 archive log compression reduces log archive storage
 - Automatically compresses log extents on-the-fly
 - Uses the same algorithm as (default) backup compression
- Simply turn it on and DB2 does the compression for you
 - Enabled by Logarchcompr1 and/or Logarchcompr2 database configuration parameter set to ON

Large SAP customer generates 60GB of log per day and they keep 8 weeks of archives

- Storing 3.3TB of archived log files
- Compression of 4x results in storage of only 825GB for 8 weeks



Client Success What Customers are Saying ...

- Customer experiences are consistent
 - Tables will compress in the range of 60% 80%
 - Indexes compress around 50%
 - Temp space will compress around 35%
 - Overall database storage savings will be between 50% and 65%



That's 65% less disk space needed to support a DB2 10 database!

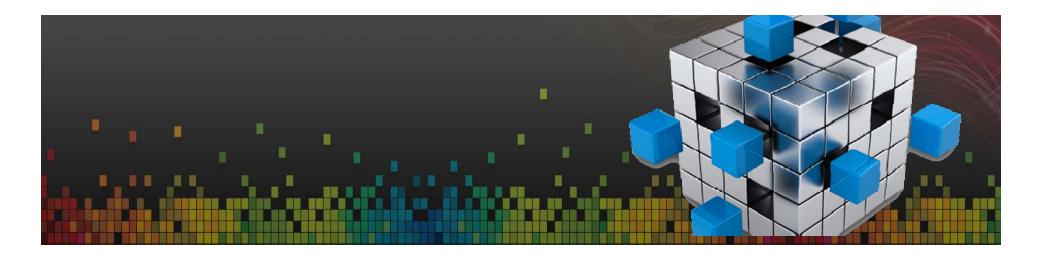


Summary

- Storage optimization through DB2 10 compression can save 55-70% of your database storage requirements
- You not only see your online database storage size shrink but often more importantly, your backup storage and disaster recovery storage is less as well
- In real customer examples storage savings are realized and performance benefits are apparent
- DB2 10 offers better compression ratios, improved ease of use and higher data availability
 - Adaptive compression provides better compression than classic row compression
 - Adaptive compression adapts to date skews over time i.e. no REORG needed to maintain compression ratio



The next steps...





The Next Steps...

- Complete the Hands on Lab for this module
 - Log onto SKI, go to "My Learning" page, and select the "In Progress" tab.
 - Find the module
 - Download the workbook and the virtual machine image
 - Follow the instructions in the workbook to complete the lab
- Complete the online quiz for this module
 - Log onto SKI, go to "My Learning" page, and select the "In Progress" tab.
 - Find the module and select the quiz
- Provide feedback on the module
 - Log onto SKI, go to "My Learning" page
 - Find the module and select the "Leave Feedback" button to leave your comments





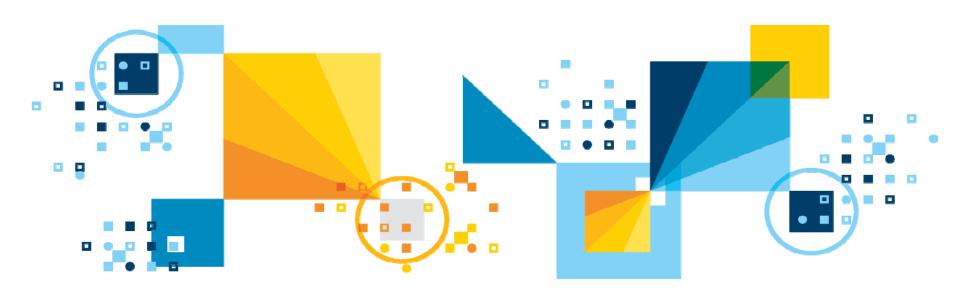
The Next Steps...

- The next set of modules to consider:
 - < Name the next set of modules that user should consider.>
- Additional Reading Material
 - Provide links etc (might be available for existing bootcamps)
 - Link 1
 - Link 2
- <Any other topics such as applicable social media links>





Questions? askdata@ca.ibm.com



January 22, 2015