

# *DB2 SQL Query Tuning with BLU Acceleration II*

Version/Revision# 1

Module ID 10304

Length 2 hours



## Table of Contents

<b>1</b>	<b>Introduction .....</b>	<b>3</b>
<b>2</b>	<b>Suggested Reading .....</b>	<b>3</b>
<b>3</b>	<b>Lab Preparation .....</b>	<b>4</b>
3.1	Start Virtual Machine and DB2.....	4
<b>4</b>	<b>Optim Query Workload Tuner Lab .....</b>	<b>6</b>
4.1	Preparing the Database.....	6
4.2	Tune a Workload .....	9
4.2.1	Launching Data Studio.....	9
4.2.2	Create a Optim Query Tuner Project.....	11
4.2.3	Run Workload Advisors .....	15
4.3	Access Path Explorer and Access Graph.....	20
4.3.1	Run Single Query Advisors.....	25
4.4	Optimization Profiles.....	29
4.4.1	Create an Optimization Profile .....	29
4.4.2	Create a New Join Sequence .....	30
4.4.3	Set Access Path to IXSCAN .....	33
4.4.4	Validate Optimization Profile .....	34
4.4.5	Compare the Access Plan change .....	34
4.4.6	Deploy Optimization Profile.....	35
<b>5</b>	<b>SQL Performance Tuning – Clean Up.....</b>	<b>37</b>
<b>6</b>	<b>Compare query processing for row-organized tables to column-organized tables lab .....</b>	<b>38</b>
6.1	Lab Preparation.....	38
6.1.1	Initialize Environment.....	38
6.2	What this exercise is about.....	38
6.2.1	Query processing for BLU and non-BLU access plans.....	38

## 1 Introduction

DB2's performance monitoring framework has been improved in DB2 10.5. Historically DB2's monitoring utilities and interfaces have been based on the monitoring elements presented by the Snapshot Monitoring stream. DB2's Health monitor, snapshot commands, administrative views and table functions all utilize the Snapshot Monitoring stream of monitoring elements.

DB2 V10 brings along with it a new and full complement of monitoring table functions and administrative views that enable you to take advantage of the monitoring framework. The monitoring framework is robust, light weight and features many identical or similar monitoring elements as the Snapshot monitoring.

This lab is intended to teach you the essentials to monitoring a DB2 database system. The scope of a "DB2 database system" extends capabilities through the Enterprise Edition, and up to and including version 10.5. This module is intended for someone who has some experience with DB2 and wishes to broaden their knowledge in the area of performance-related monitoring.

These labs are performed on a VMware image of a 64bit SUSE Linux machine with minimal resources declared. There are some commands used in the lab that are typically found on Linux/Unix systems, however basic Unix command-line tools can be downloaded and easily employed on Windows-based machines at no cost. Additionally, comparable commands are available in the Windows environment, like find in place of grep.

### **There are two main tips that will make the lab progress better.**

1. Use the up arrow key to recall previous commands. These commands can be edited and re-executed as needed.
2. Follow all the instructions, don't skip steps. These labs have been created so that they can be rerun but to avoid that please follow all instructions.

## 2 Suggested Reading

### **Information Management Technical Library at developerWorks**

Contains articles and tutorials and best practices

<http://www.ibm.com/developerworks/views/db2/library.jsp>

<http://www.ibm.com/developerworks/data/bestpractices>

### **An Expert's Guide to DB2 Technology**

Great read and useful information.

<http://blogs.ittoolbox.com/database/technology>

### **DB2 10 Fundamentals Certification Study Guide**

Learn the basics and get ready for certification

### 3 Lab Preparation

Before getting started with the lab exercises you will set up your environment then work through a series of exercises focused on DB2 performance monitoring techniques.

#### 3.1 Start Virtual Machine and DB2

You may skip this section if you have started your VM and db2 instance in a previous lab or exercise.

1. Choose the **DB2\_BLU-AWSE\_10.5....vmx** file when first opening the lab image in the **/DB2\_BLU\_PerfMonitoring...** folder and run it.



**Figure 1 – Opening a VMware Image File**

2. If this is the first time this virtual machine has been started, a series of licensing agreements will prompt you. To proceed, acknowledge by responding by selecting the “**Yes, I Agree to the License Agreement**” selection.



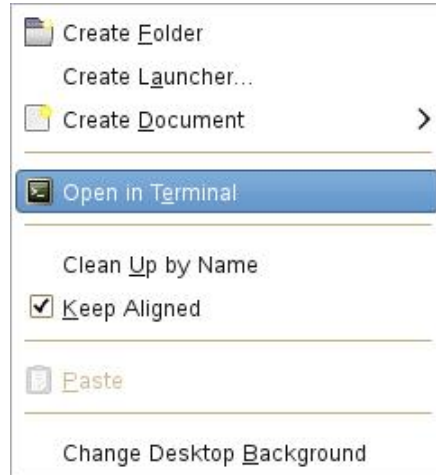
Figure 2 – License Agreements

3. Login in at the console prompt as **db2inst1**. You will be prompted for the password which is **password**. You may be prompted for the root **password** as well, if so, it is also password.



Figure 3 – Login credentials

4. Open a terminal window as follows:
  - Open a new terminal window by right-clicking on the **Desktop** and choosing the “**Open Terminal**” item:



**Figure 4 – Opening a Terminal**

From the new terminal session start the Database Manager as follows:

- Enter the command “**db2start**” to start the database instance. Make sure the instance is started successfully. (**Note:** If you notice a message that says “The database manager is already active”, that’s “OK. Please continue to the next step.)
- In order to enable the aliases defined for this lab, copy the /home/db2inst1/labbashrc to /home/db2inst1/.bashrc and source that file.

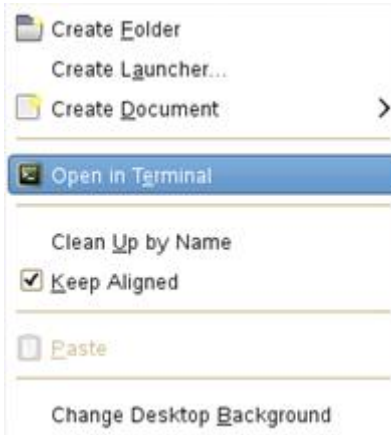
```
cd ~  
cp labbashrc .bashrc  
. .bashrc
```

## 4 Optim Query Workload Tuner Lab

In this part of the lab, you will work with the Optim Query Tuner graphical tool with DB2 10. You will run a workload on the database server and use advisors to obtain recommendations indexes, statistics and views are suggested to improve the workload performance. The recommendations received will be applied to the database server and the workload rerun.

### 4.1 Preparing the Database

1. Open a new terminal window by right-clicking on the **Desktop** and choosing the “**Open Terminal**” item:



**Figure 2 – Opening a Terminal**

2. Use the terminal window directory to prepare the database with script oqwt01.sh located in /home/db2inst1/Documents/LabScripts/db2pt/querytuning/oqwt. Run the following command:

```
oqwtDir  
./oqwt01.sh
```

After some time, the following should display, indicating that everything has been processed successfully:

```
-----  
Run EXPLAIN.DDL in the SAMPLE database  
-----
```

```
Database Connection Information
```

```
Database server      = DB2/LINUX8664 10.5.3  
SQL authorization ID = DB2INST1  
Local database alias = SAMPLE
```

```
DB20000I The SQL command completed successfully.  
-----
```

```
Applying OQWT License to the SAMPLE database  
-----
```

```
Database Connection Information
```

```
Database server      = DB2/LINUX8664 10.5.3  
SQL authorization ID = DB2INST1  
Local database alias = SAMPLE
```

```
DROP FUNCTION DB2OE.QT_LIC
```

```
DB20000I The SQL command completed successfully.
```

```
CREATE FUNCTION DB2OE.QT_LIC() RETURNS VARCHAR(255) LANGUAGE SQL CONTAINS SQL NO EXTERNAL ACTION  
DETERMINISTIC RETURN VARCHAR('wVLwF/a4OBBLTyNX4bsUKw==')
```

```
DB20000I The SQL command completed successfully.
```

```
GRANT EXECUTE ON FUNCTION DB2OE.QT_LIC TO PUBLIC WITH GRANT OPTION
```

```
DB20000I The SQL command completed successfully.
```

```
DB20000I The SQL command completed successfully.  
-----
```

```
Note: You have applied OQWT license to the SAMPLE database
```



## 4.2 Tune a Workload

### 4.2.1 Launching Data Studio

1. Open IBM Data Studio by double-clicking on the IBM Data Studio icon on the Desktop.



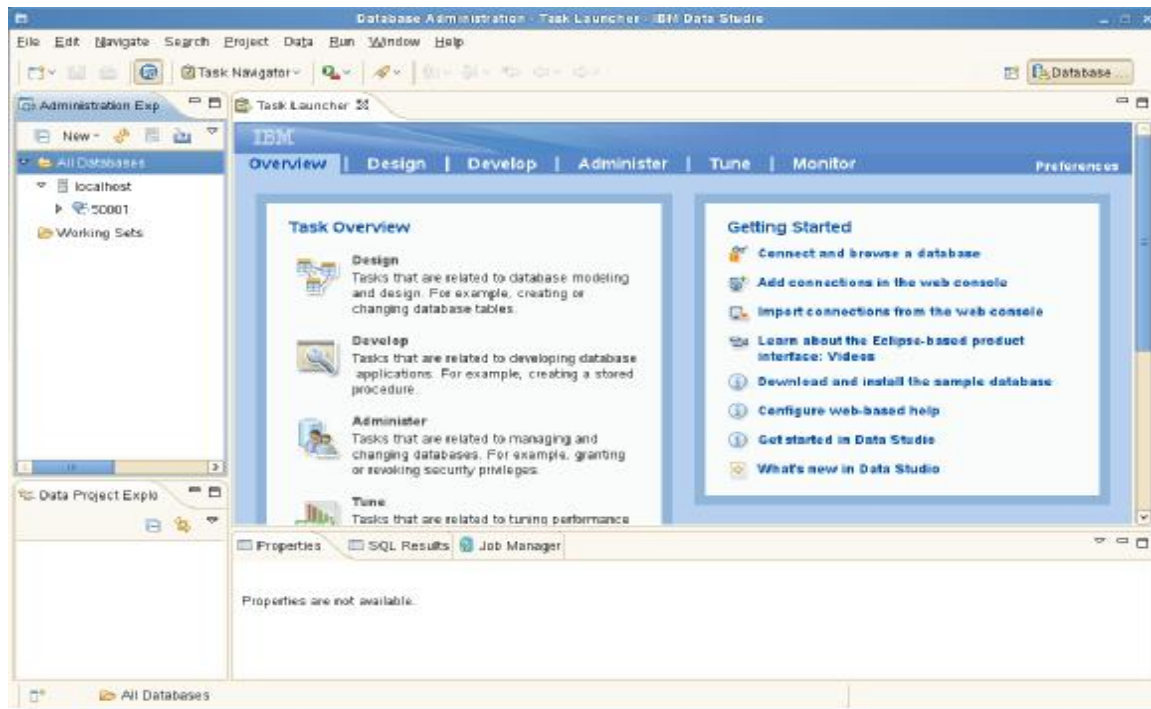
**Figure 3 – Open IBM Data Studio**

2. If this is the first time Data Studio is launched, it may take a moment. Click **OK** to accept the default workspace and check the option “**Use this as the default and do not ask again**”.



**Figure 4 – Select a workspace**

3. By default, the Database Administration perspective is opened. Your workspace should look like the following:



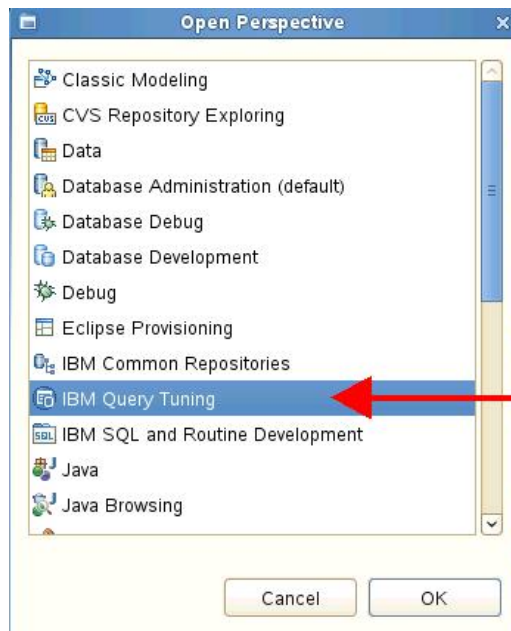
**Figure 5 – Database Administration perspective**

4. We want to change to the **IBM Query Tuning** perspective. To do this, click on the “**Open Perspective**” icon and click **Other**, as shown below.



**Figure 6 – Open perspective**

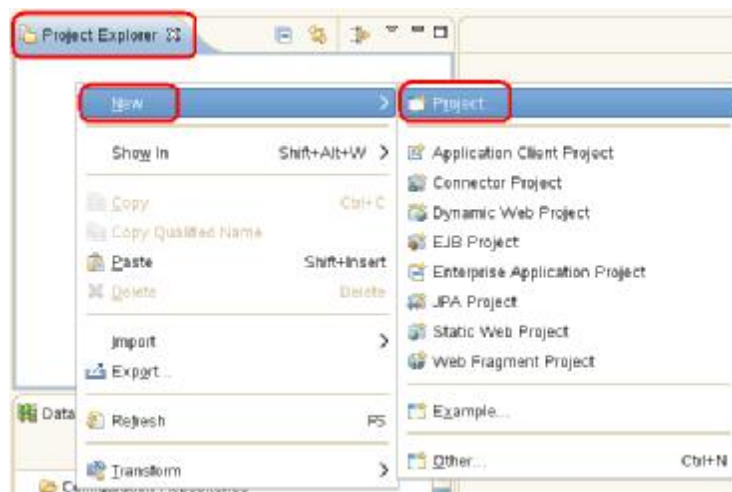
Click on IBM Query Tuning and click on the **OK** Button.



**Figure 7 – IBM Query Tuning Perspective**

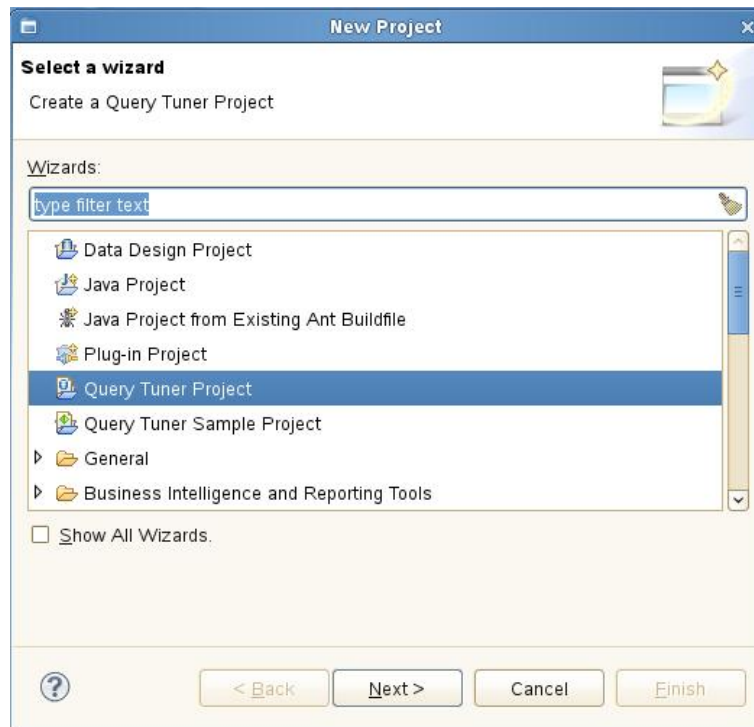
#### 4.2.2 Create a Optim Query Tuner Project

1. Move your mouse pointer to the **Project Explorer** section, right-click and select **New**, and then **Project...**



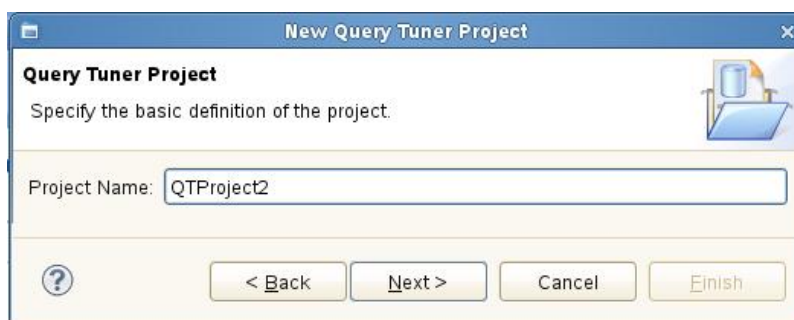
**Figure 8 – Create New Project**

2. Select **Query Tuner Project** from the wizard list and click **Next**.



**Figure 9 – Query Tuner Project Wizard**

3. Leave the **Project Name** as suggested in the wizard and click **Next**.



**Figure 10 – Query Tuner Project Name**

4. Select the **SAMPLE** database for the connection and click **Next**. If necessary create connection following steps 6 to 8 in Lab 3.3 of last module's workbook. Remember use SAMPLE as database name. Then click **Finish**.

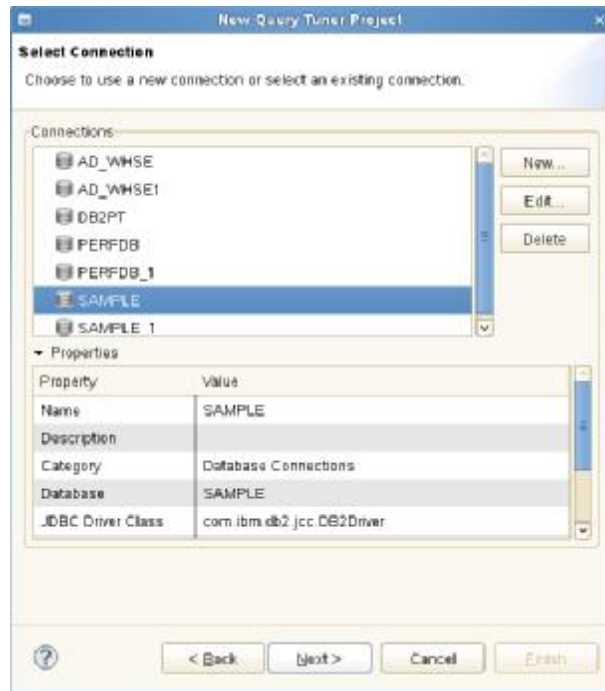


Figure 11 – Select Database connection

5. In this lab, we will use a workload file, with multiple SQL statements. First select **File** and click **Browse**.


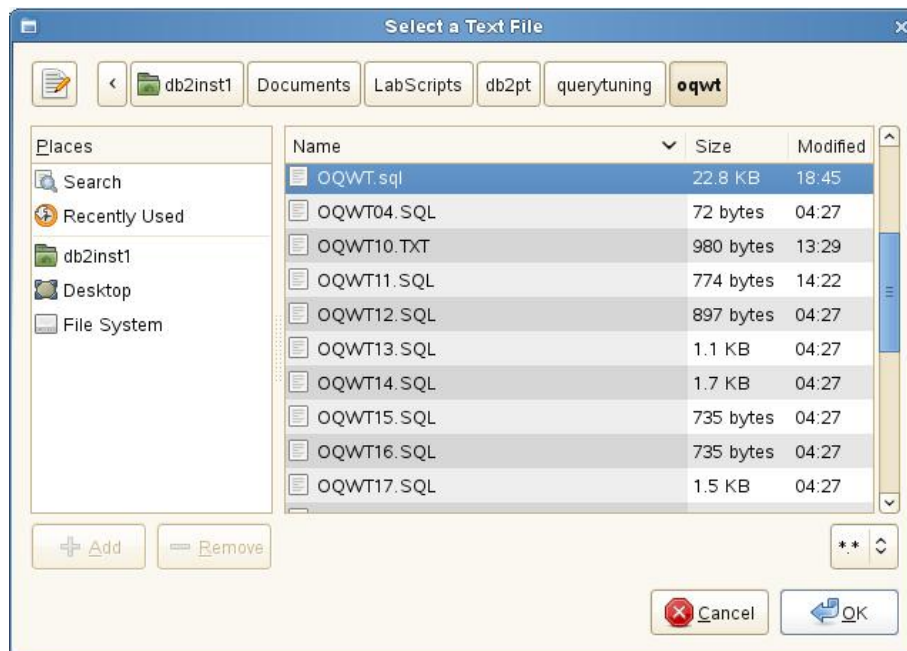


Figure 12 – Capture SQL Sentences

6. Select the file **OQWT.sql** located in `/home/db2inst1/Documents/LabScripts/db2pt/querytuning/oqwt`, and click **OK**.



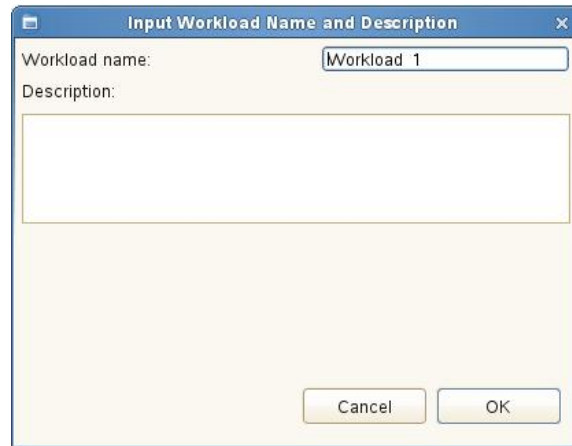
**Figure 13 – Select query workload file**

7. Use “;” as statement delimiter, click **Capture**. Once the statement appears on screen, click **Save All to Workload...**



**Figure 14 – Capture workload**

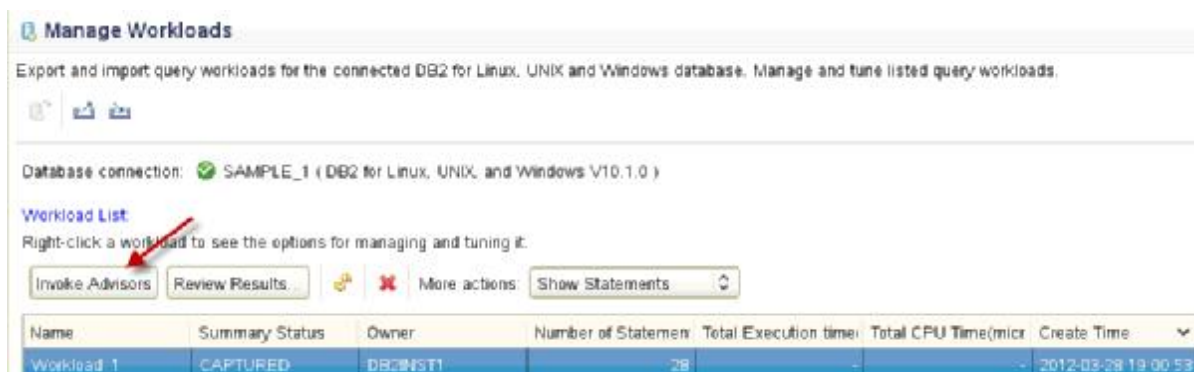
8. Accept the default workload name and click **OK**.



**Figure 15 – Workload Name**

#### 4.2.3 Run Workload Advisors

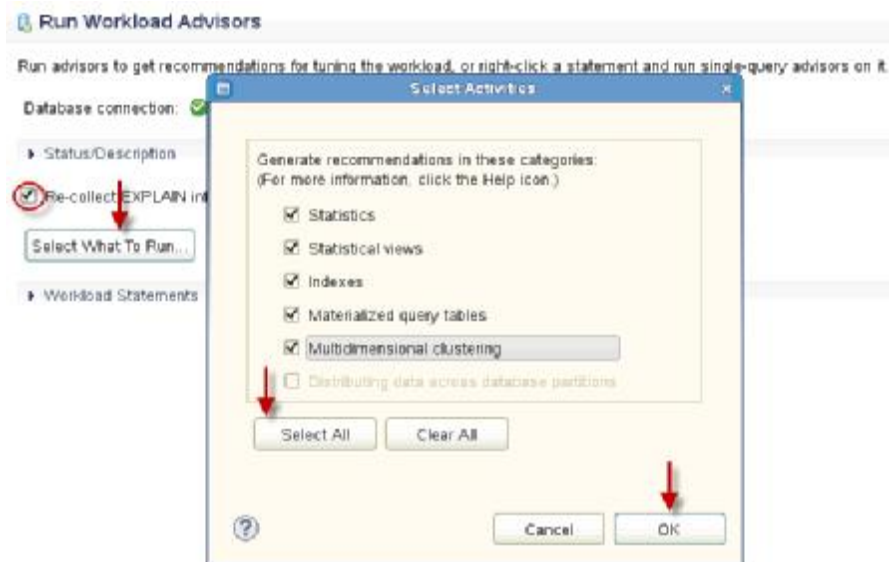
1. Click **Invoke Advisors** to call workload advisors and tools.



**Figure 16 – Invoke Advisors**

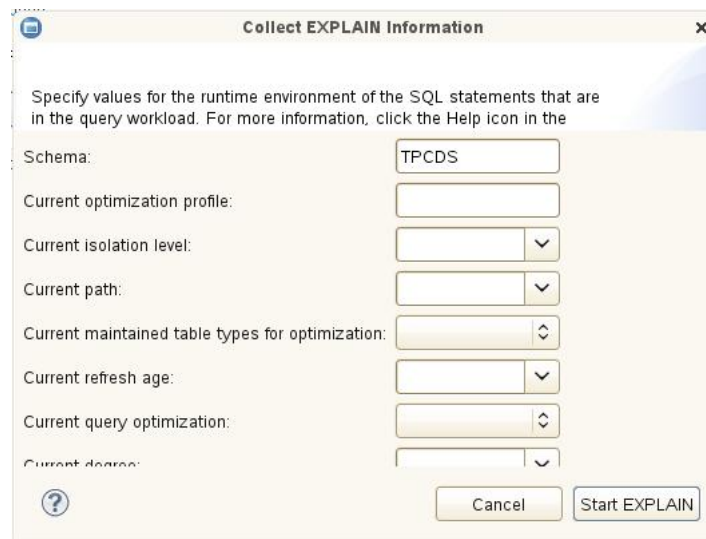
2. Check **Re-collect EXPLAIN information** before running workload advisors. Click **Select What To Run....** Click **Select All**, and click **OK**.





**Figure 17 – Select categories for recommendation**

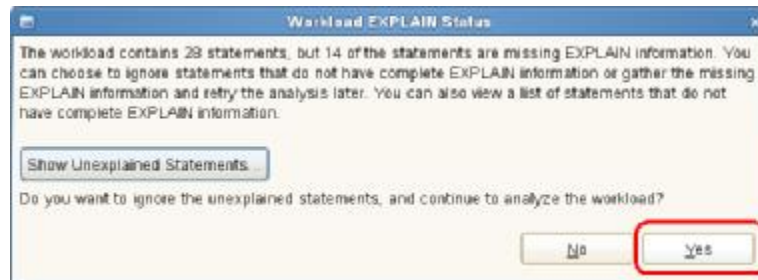
3. Specify the **TPCDS** schema and click **Start EXPLAIN**.



**Figure 18 – Specify Explain Values**

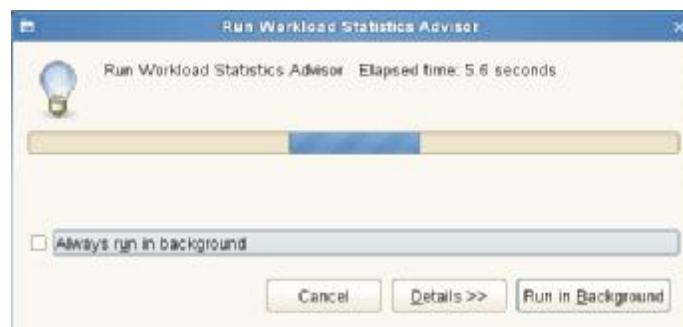
4. There may be Unexplained Statements such as “set current schema= ‘TPCDS’”, Click **Yes** to continue.





**Figure 19 – Status**

5. Workload advisors will start working on the workload. Please wait for this to finish.



**Figure 20 – Running advisors**

6. In case if an error window is displayed, press OK.
7. After this, the recommendations are presented. We'll review them in the next few steps.

Status/Description				
Statements Summary Statistics Indexes Statistical Views				
Item Analyzed	Result	Recommendation Started	Recommendation Completed	
Statistics	New recommendations were generated.	2012-03-28 19:12:11	2012-03-28 19:12:24	
Indexes	New recommendations were generated.	2012-03-28 19:13:10	2012-03-28 19:13:00	
Statistical Views	New recommendations were generated.	2012-03-28 19:12:27	2012-03-28 19:13:01	

**Figure 21 – Recommendation Summary**

8. Click the **Statistics** tab and click **View RUNSTATS**.



10. Click **Finish** to execute the RUNSTATS COMMANDS on the database. Click **OK** in the confirmation window.



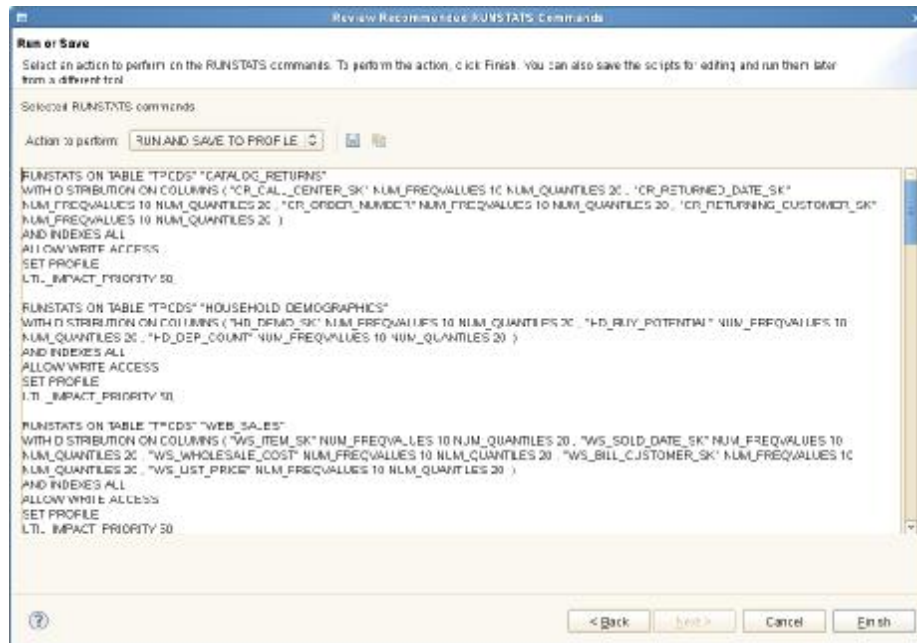


Figure 24 – RUNSTATS Commands

**Note:** In DB2 10, a number of improvements have been made to the RUNSTATS command to make statistics gathering faster. The command parameters have also been simplified.

- Click the **Indexes** tab to explore recommendations for indexes.

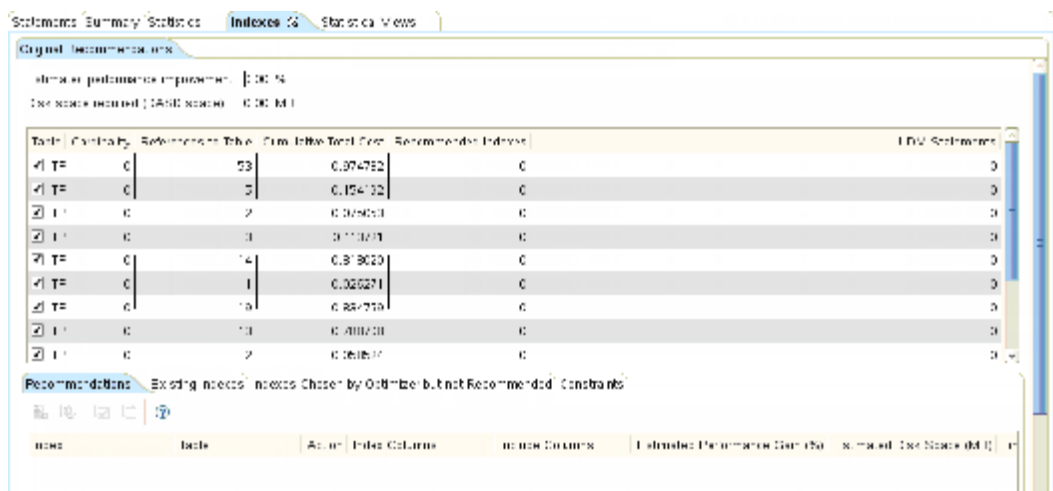


Figure 25 – Indexes Recommendations

- Click **Statistical Views** tab to see the recommendations. The use of statistical view helps to improve the performance of queries and OQWT helps to determine these views based upon the workload. It is not an easy task to do by hand.

Highlight a statistical view to see details about it below the summary table. To view the scripts for creating or modifying statistical views, highlight the table.

Estimated RUNSTATS execution time: 0.00000 minutes Recommended new statistical views: 0  
Recommendations for creating statistical views: 0 Existing statistical views without recommendations: 0

View Script

Checks to View Scripts	Owner	Name	Status	Tables in Definition	Affected Statements	Estimated Execution Time (minutes)	Estimated Rows
<input checked="" type="checkbox"/>	FCDS	SVIEW01	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW02	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW03	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW04	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW05	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW06	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW07	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW08	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW09	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW10	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW11	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW12	PENDING	2	1	0.000000	
<input checked="" type="checkbox"/>	FCDS	SVIEW13	PENDING	2	1	0.000000	

Database: DB2LUW1

Statements Accessed by: SVIEW01

Statement	Number of executions	Estimated Execution Time (minutes)
select ...	1	0.00000

Figure 26 – Statistical Views Recommendations

13. Close project tab as shown, **Save** changes and **Exit**.

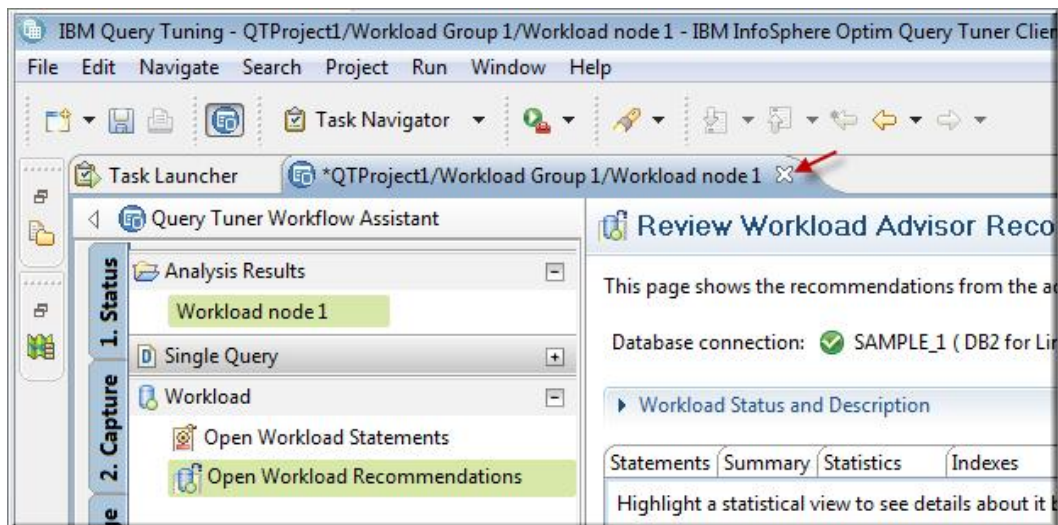
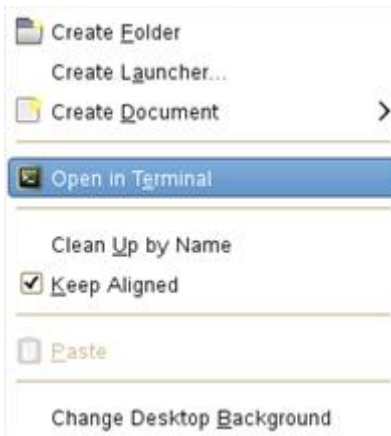


Figure 27 – Close Project Tab

## 4.3 Access Path Explorer and Access Graph

1. Open a new terminal window by right-clicking on the **Desktop** and choosing “**Open in Terminal**.”:

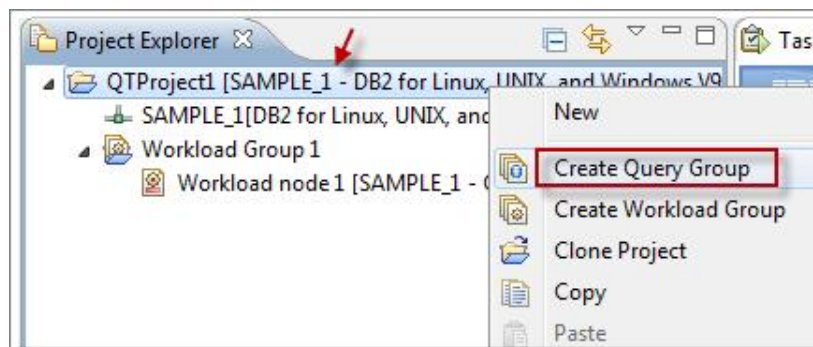


**Figure 28 – Opening a Terminal**

2. Run the oqwt03.sh script located in directory /home/db2inst1/Documents/LabScripts/db2pt/querytuning/oqwt to create two tables, wait for a message stating that it “Successfully Executed”.

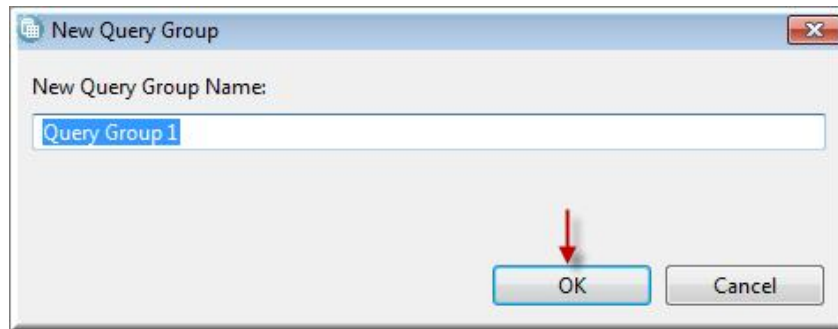
```
oqwt_dir  
./oqwt03.sh
```

3. Back in Data Studio, right-click on the QTPProject1 and select **Create Query Group**. [We saved this project in the previous section.]



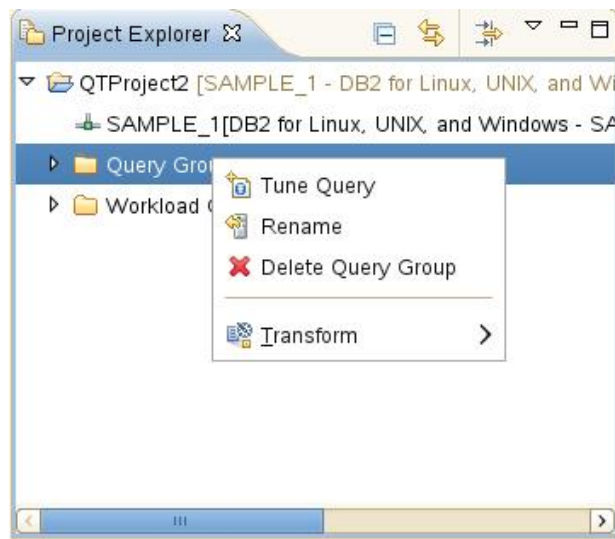
**Figure 29 – Create Query Group**

4. Accept the default name and click **OK**.



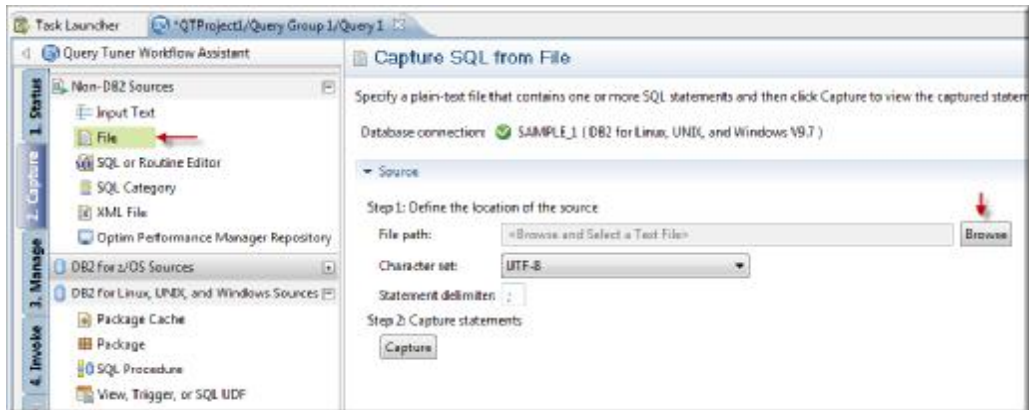
**Figure 30 – Default Query Group name**

5. Right click on Query Group X, where X is the generated group number, and select the option **Tune Query**.



**Figure 31 – Tune Query**

6. As in the previous section we need a file with a SQL statement to optimize Click **File**, and then click **Browse**.



**Figure 32 – Select file with SQL**



7. Select the file **OQWT04.sql** located in `/home/db2inst1/Documents/LabScripts/db2pt/querytuning/oqwt`. Click **OK**, use “;” as statement delimiter and click **Capture**. Once the statement appears on screen, click **Save All to Workload...**

▼ Source

Step 1: Define the location of the source

Qualifier:

File path:

Character set:

Statement delimiter:

Step 2: Capture statements

▼ Captured Statements

Step 3: Select a statement and click on Invoke Advisors and Tools, or save all statements to a new workload.

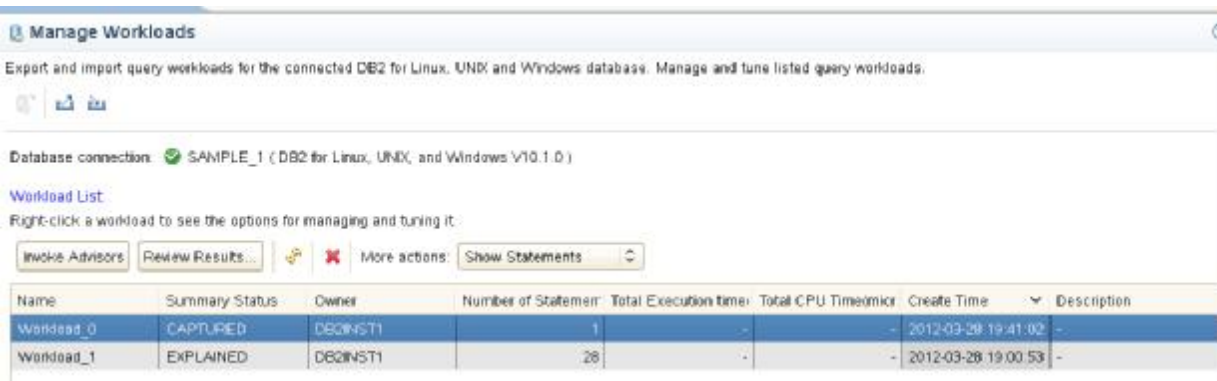
The number of captured statements is 1.

Number	Statement Text
0	SELECT * from OQWT_T1 T1, OQWT_T2 T2 WHERE T1.c1 = T2.c2 and T1.c3 >100

**Figure 33 – Capture SQL from a File**



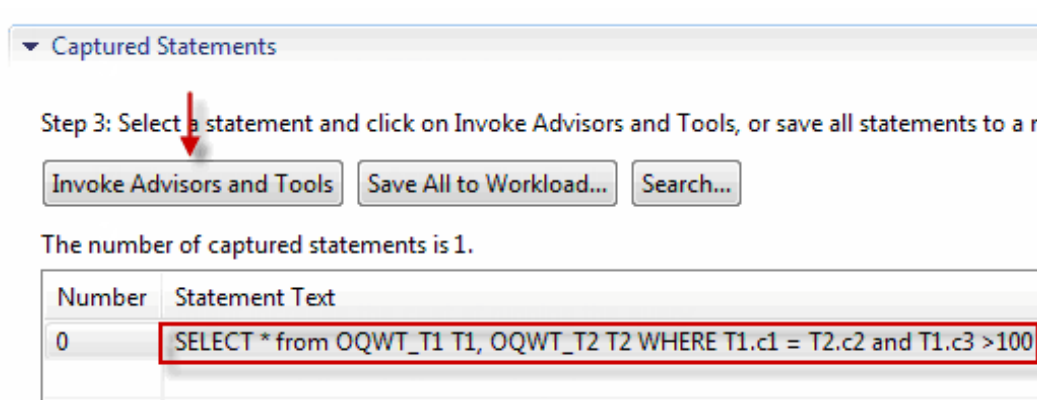
8. Select workload name you just saved in the previous step, double click to open.



**Figure 34 – Manage Workload**

#### 4.3.1 Run Single Query Advisors

1. Click the vertical tab “2. Capture” so that we can run advisors. Click **Invoke Advisors and Tools**.



**Figure 35 – Invoke Advisors**

2. Click **Select What To Run...**

Schema: **DB2INST1** ←

☐ Re-EXPLAIN the query

▶ EXPLAIN options and runtime environment options

Select What To Run... Run SQL

▼ Query Text - Query 3

```
SELECT * from OQWT_T1 T1, OQWT_T2 T2 WHERE T1.c1 = T2.c2 and T1.c3 >100
```

**Figure 36 – Invoke Advisors**

3. Click **Select All**, then **OK**.

Select Activities

Analysis tools:

- ☒ Format and annotate SQL statement
- ☒ Display access plan graph
- ☒ Show access plan in Access Plan Explorer
- ☒ Collect actual execution values if the SQL is a SELECT statement  
Note: This option is for SELECT statements;  
It might increase the cost of running the query.

Generate recommendations in these categories:  
(For more information, click the Help icon.)

- ☒ Statistics
- ☒ Query revision
- ☒ Access path
- ☒ Indexes

Generate reports:

- ☒ Recommendation summary

Select All Select Defaults Clear All

☐ Save changed query  
Name of query: \_\_\_\_\_

☐ Save new analysis result  
Name of analysis result: \_\_\_\_\_

Cancel OK

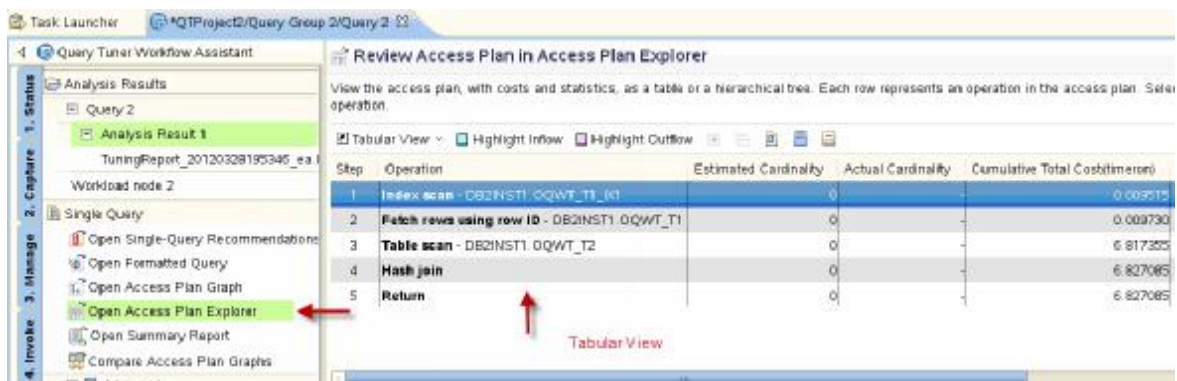
**Figure 37 – Select Activities**

4. The single query tune advisor will start to run. Please wait for this to finish.



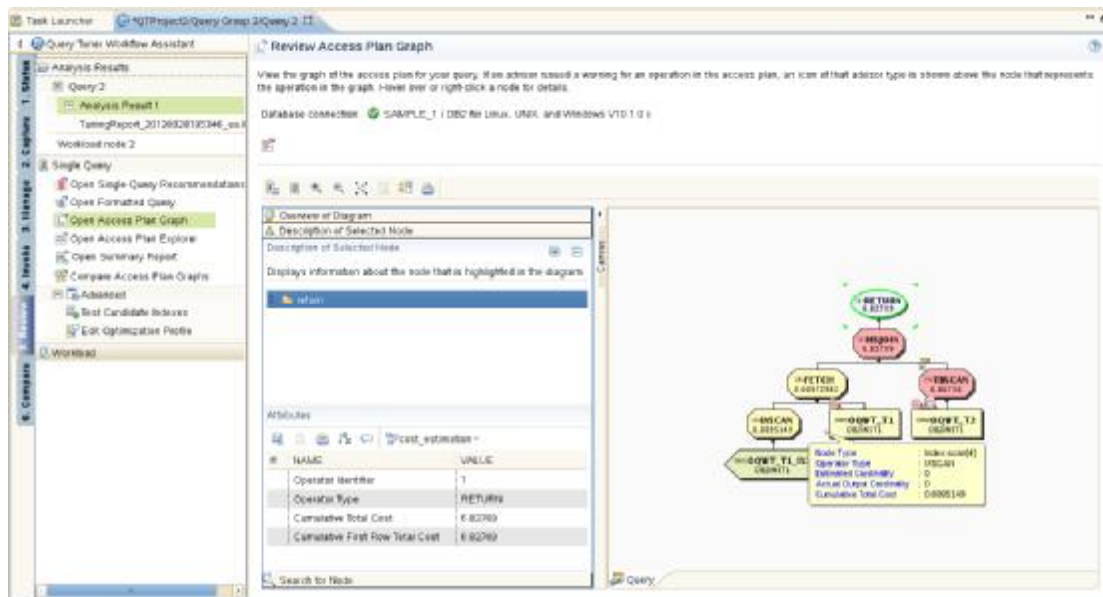
**Figure 38 –Tuning Query**

5. Click **Open Access Plan Explorer**.



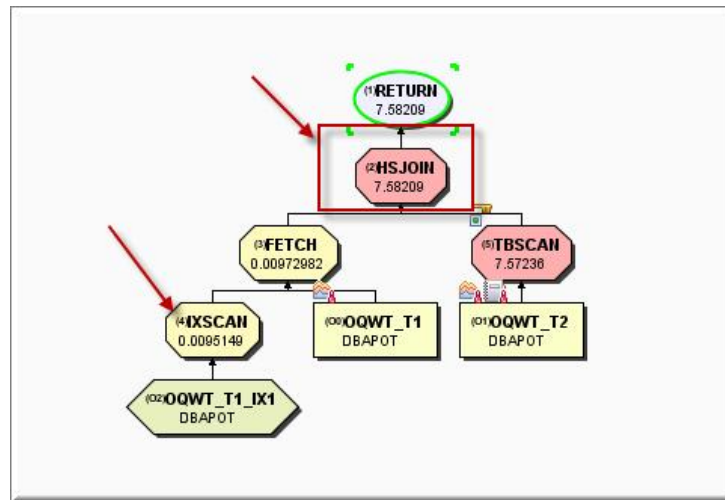
**Figure 39 – Access Plan Tabular View**

6. Click **Open Access Plan Graph**.



**Figure 40 – Access Plan Graph View**

- Notice in the access plan graph, there is a HASH JOIN between a full table scan on the table OQWT\_T2 and IXSCAN on the table OQWT\_T1. We will take this as an example and show you how to create an optimization profile to influence the query access path.



**Figure 41 – Access Plan Graph Detail**

## 4.4 Optimization Profiles

An optimization profile is a DB2 feature that allows a DBA to provide a “hint” for an SQL statement. Normally, the DB2 optimizer does not require “hints”, but it may be necessary to influence the optimizer to do things in a particular fashion ONLY for a particular query.

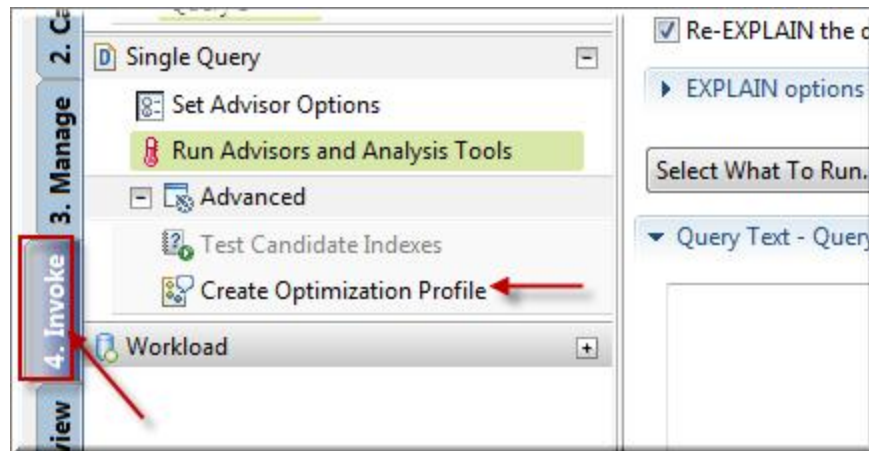
To set up the optimization profile, we will use the access plan graph of the query we saw in the previous section.

We will create an optimization profile using the following requirements.

- OQWT\_T1 table to be accessed with IXSCAN and it should use index OQWT\_T1\_IK1
- OQWT\_T2 must be a leading table in the join sequence and the join method for OQWT\_T1 must be nested-loop join and not the hash-join shown in the access plan graph

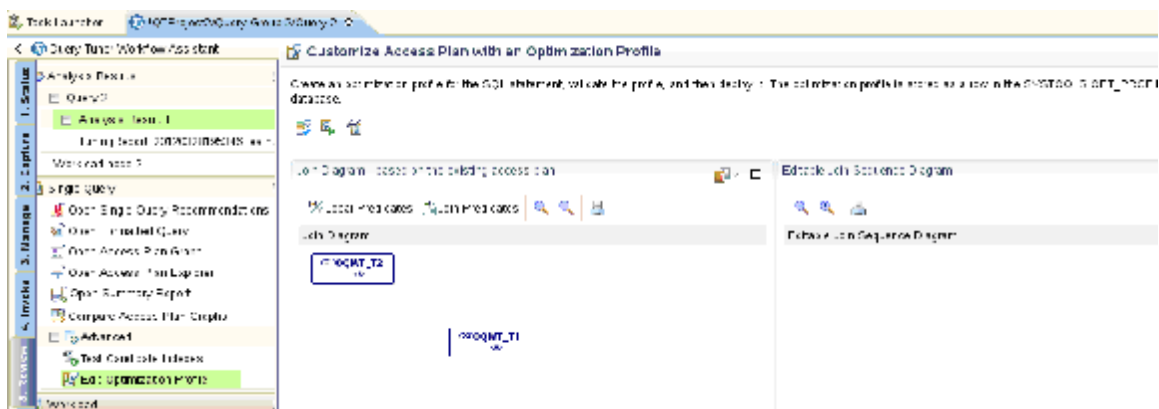
### 4.4.1 Create an Optimization Profile

- Click vertical tab “4. Invoke.” Click **Create Optimization Profile.**



**Figure 42 – Create Optimization Profile**

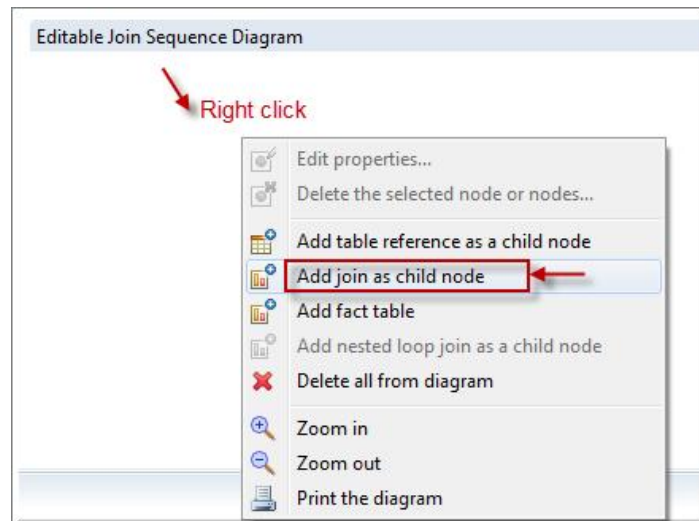
- Please notice that the Plan-Hint 1 is created. You will see two panes on the right-hand side. One of them is the Join Diagram, based on existing access plan. The other is the Editable Join Sequence Diagram. We will work from the second graph window.



**Figure 43 – Create Plan Hint**

#### 4.4.2 Create a New Join Sequence

- Right-click anywhere in the Editable Join Sequence Diagram. Select **Add Join as Child node**.



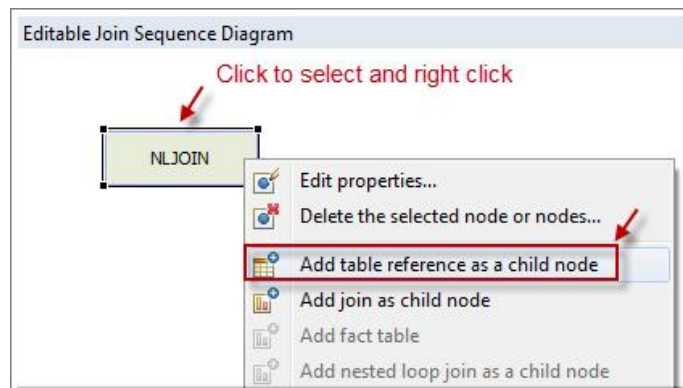
**Figure 44 – Adding Join**

2. Click the **Join Type** dropdown, select **NLJOIN** and click **OK**.



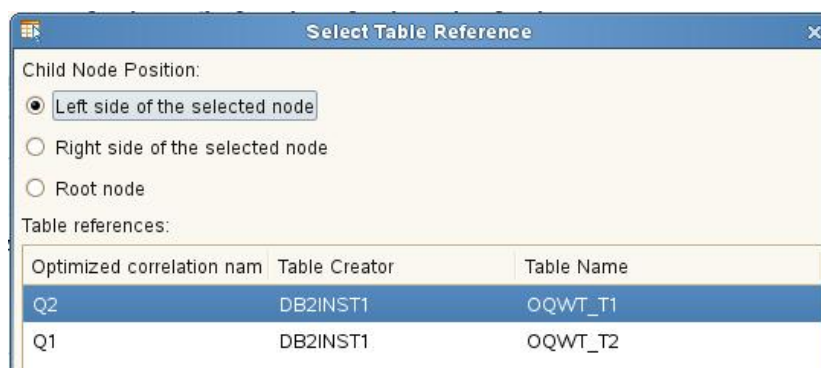
**Figure 45 – Join Type**

3. An NLJOIN box is created in the diagram window. Click to select it. Right-click and select **Add table reference as a child node**.



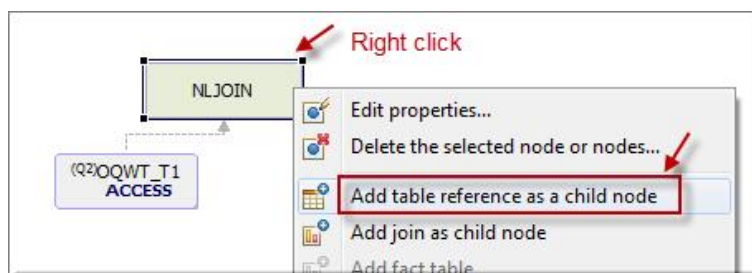
**Figure 46 – Add table reference**

4. Select Q2 (OQWT\_T1) and click **OK**.



**Figure 47 – Select Right Table Reference**

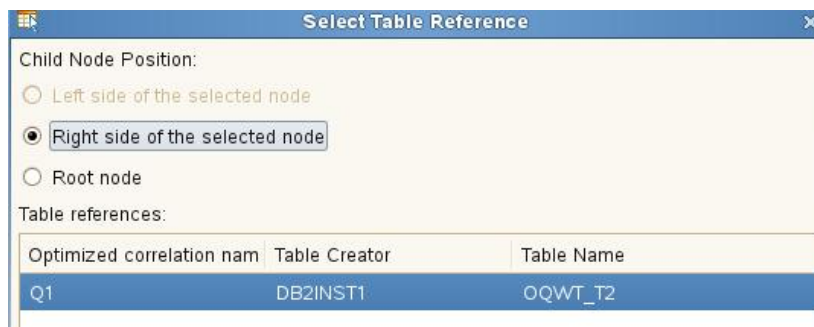
5. Right-click on the NLJOIN box. Select **Add table reference as a child node**.



**Figure 48 – Add Table Reference**

6. Click Q1 line (OQWT\_T2). Click **OK** in the dialog box.

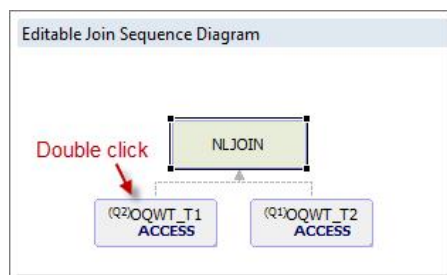




**Figure 49 – Select left Table Reference**

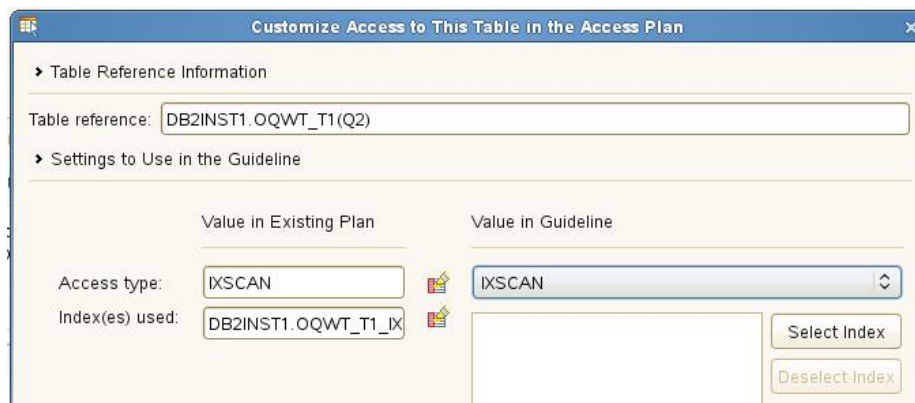
#### 4.4.3 Set Access Path to IXSCAN

1. You will notice the nested loop join sequence created between two tables as shown. Double click the OQWT\_T1 box.



**Figure 50 – Join Sequence Diagram**

2. Click the **Value in Guideline** dropdown. Select **IXSCAN** and click **Select Index**. Check DB2INST1.OQWT\_T1\_IX1 and click **OK**.

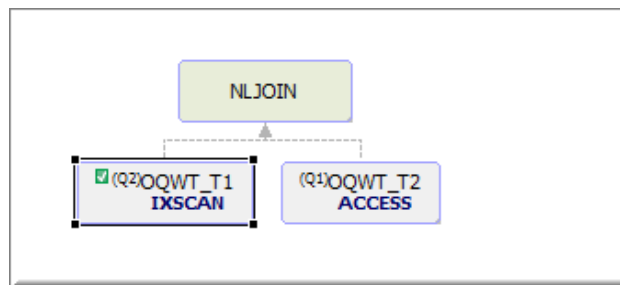


**Figure 51 – Select access type and Index**



**Figure 52 – Select DB2INST1.OQWT\_T1\_IX1Index**


3. You will notice the following join graph between two tables.



**Figure 53 – Join Graph**

#### 4.4.4 Validate Optimization Profile

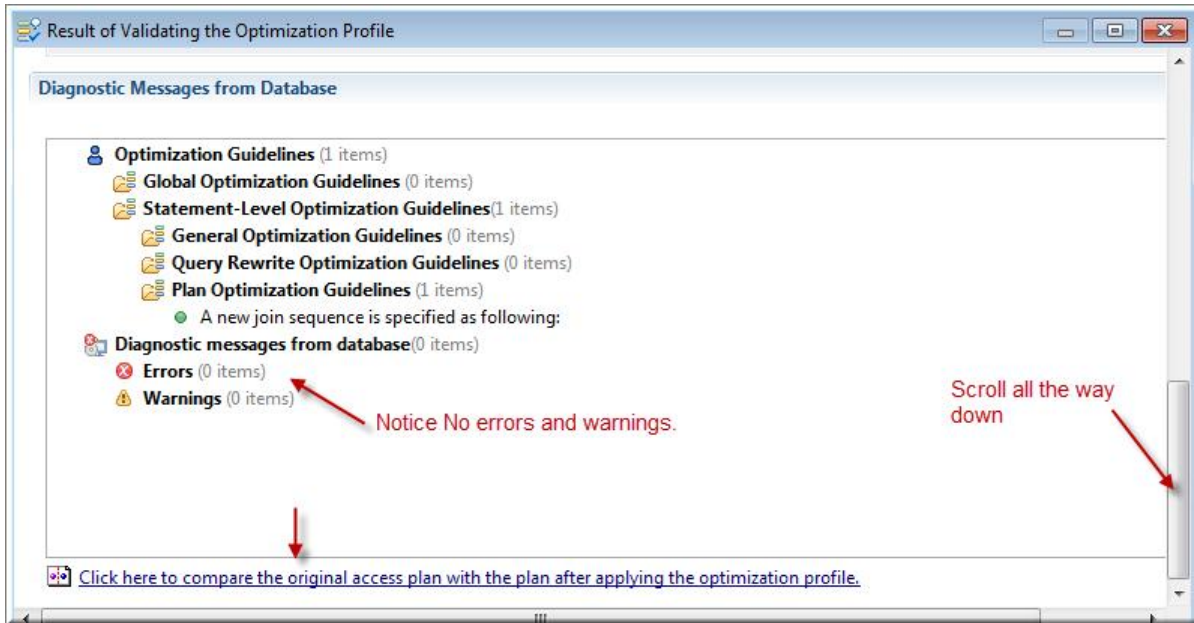
1. Click the **Validate Optimization Profile** icon, which is the left one and accept all default values. Click **Validate**. Please wait for the validation process to complete.



**Figure 54 – Validate Optimization profile**

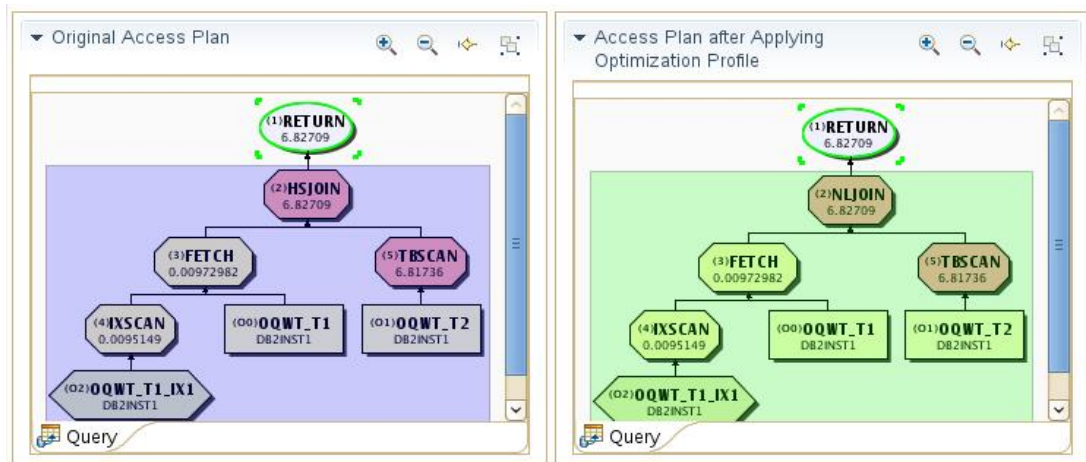
#### 4.4.5 Compare the Access Plan change

1. After validation, a new window opens with the validation results. Scroll all the way down. Notice that there are no errors or warnings.
2. Click the link (at the bottom of the window) to compare the new access path after applying the optimization profile.



**Figure 55 – Validation Results**

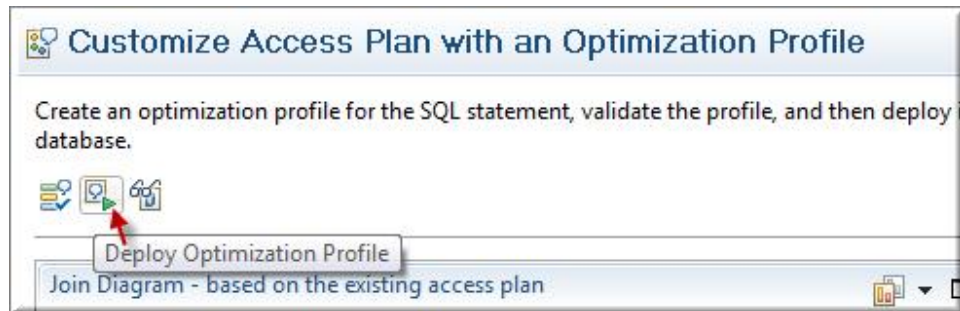
3. A new window will open with two access path graphs. Compare them and notice that the new plan is using the nested loop join. Close the window.



**Figure 56 – Access Plan comparison**

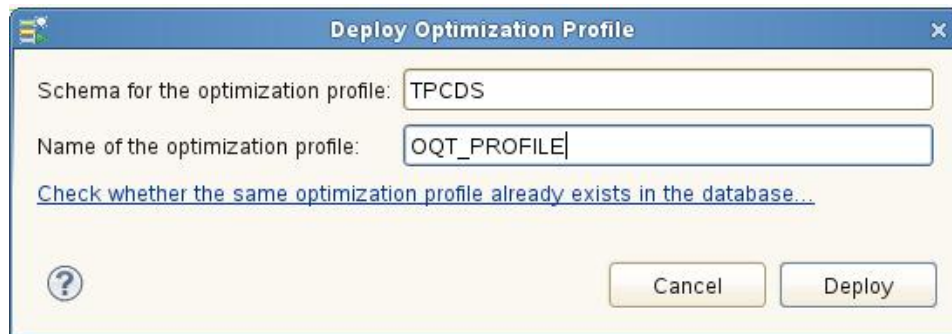
#### 4.4.6 Deploy Optimization Profile

1. Click the second icon, which is Deploy Optimization Profile.



**Figure 57 – Deploy Optimization Profile**

2. Enter **OQT\_PROFILE** as the profile name and click **Deploy**.



**Figure 58 – Schema and Optimization Profile Name**

3. A new window will open. Scroll to the bottom, browsing through the script generated. Click **OK** to deploy this optimization profile to the database.

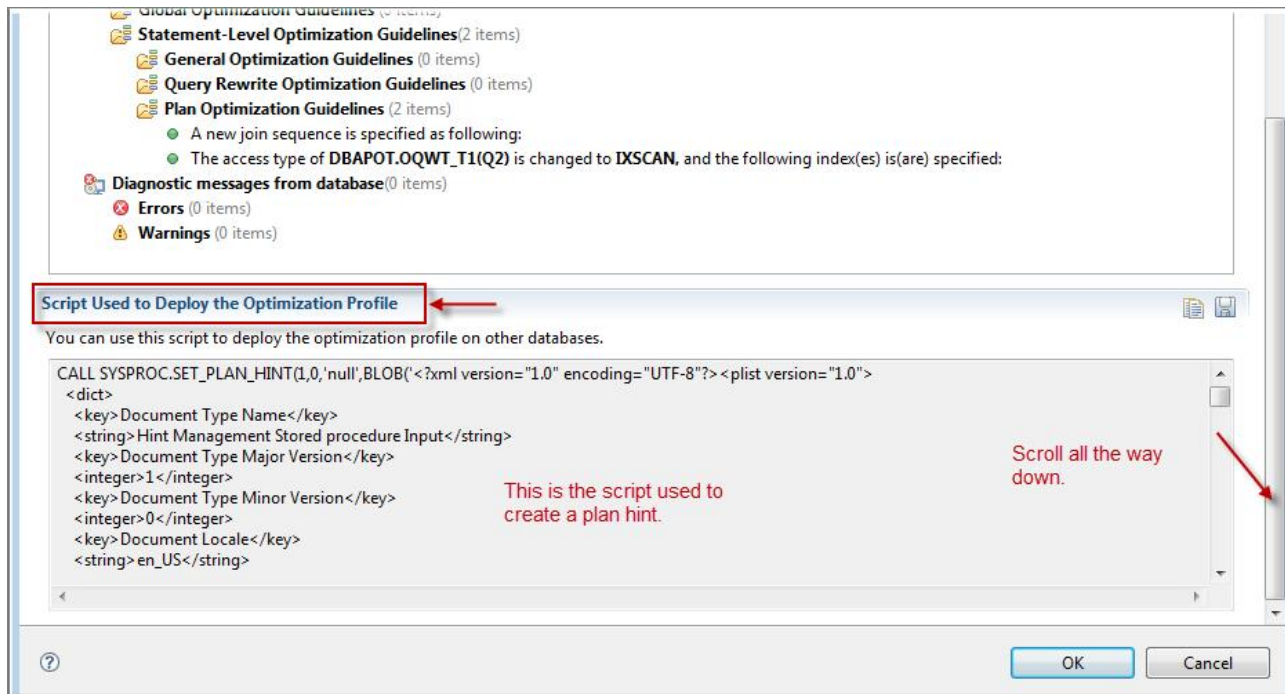


Figure 59 – Deploy

4. Close QTPProject1 as shown below. Click **Yes** to save changes.

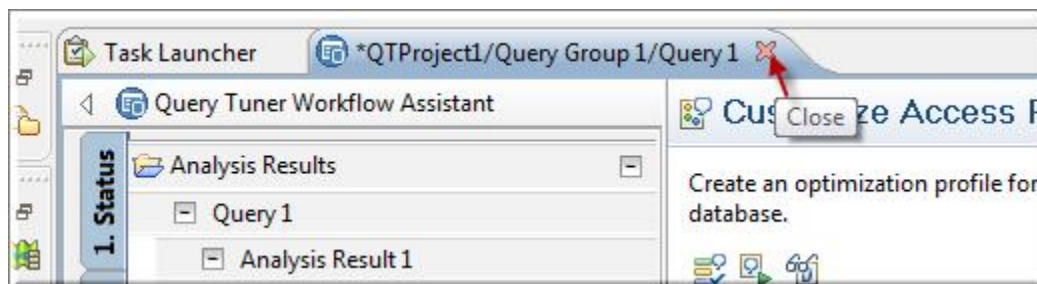


Figure 60 – Close Tab

## 5 SQL Performance Tuning – Clean Up

Run the following script to reset your lab environment.

```
qtdir
. cleanup.sh
```

## 6 Compare query processing for row-organized tables to column-organized tables lab

### 6.1 Lab Preparation

Before getting started with the lab exercises you will set up your environment then work through a series of exercises focused on DB2 performance monitoring techniques.

#### 6.1.1 Initialize Environment

You have been provided with a script, which creates a database, the needed tablespaces and the tables that you will work with in this portion of the lab. **These are mandatory steps and must be completed, as mandatory steps so often are.**

1. If you no longer have the terminal opened from previous exercise: By right clicking on the Gnome desktop (as in 3.1) **open a new terminal window.**
2. While logged in as user “**db2inst1**”, change your working directory to the folder containing the scripts.

```
/home/db2inst1/Documents/LabScripts/db2pt/db2blu
```

### 6.2 What this exercise is about

In this part of the BLU Acceleration exercise, we compare various performance characteristics of SQL queries using the DB2 BLU Acceleration with the same queries based on row-organized tables with standard indexes. We will look at the access plans using the **db2exfmt** explain tool to review the pre-execution estimated processing cost.

#### 6.2.1 Query processing for BLU and non-BLU access plans

1. In first SQL query that we will analyze accesses a single table and produces a summarized result from a subset of the table data. We will use the **db2batch** application to execute the query and report basic elapsed time statistics. We will also include several SQL queries in the db2batch sequence that return important DB2 performance metrics that we will use to compare the row-organized and columns-organized table processing.

We will start by running the SQL query using the row-organized tables. The file *rowquery1.sql* contains the following statements:

```
set current explain mode yes;
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as
br_trans
  FROM ROWORG.HISTORY AS HISTORY
 WHERE HISTORY.BRANCH_ID between 10 and 20
 GROUP BY HISTORY.BRANCH_ID
 ORDER BY HISTORY.BRANCH_ID ASC ;
set current explain mode no;
```

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,  
      ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as  
total_l_reads  
      from table(mon_get_connection(null,-1)) as con  
      where application_name = 'db2batch'  
      ;  
  
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time,  
total_wait_time  
      from table(mon_get_connection(null,-1)) as con  
      where application_name = 'db2batch'  
      ;
```

We will define several NOT ENFORCED constraints on the column-organized tables to provide the optimizer information similar to the unique indexes that are defined for the row-organized tables.

We will collect new table statistics for the column-organized tables

The CURRENT EXPLAIN setting will generate the explain data that can be formatted using **db2exfmt** to show estimated processing costs. We will use the file *explain.ddl* to create a new set of explain tables.

From your Linux VMware, login with the userid **db2inst1**. Start a new terminal session. In the terminal session, issued the following commands:

```
db2 force application all  
db2 terminate  
db2 activate db db2blu  
db2 -tvf explain.ddl  
db2 -tvf pkeys.ddl
```

The output should look similar to the following:

```
alter table colorg.acct add primary key (acct_id) not enforced  
DB20000I  The SQL command completed successfully.  
  
alter table colorg.branch add primary key (branch_id) not enforced  
DB20000I  The SQL command completed successfully.  
  
alter table colorg.teller add primary key (teller_id) not enforced  
DB20000I  The SQL command completed successfully.  
  
alter table colorg.history add constraint fkeyteller foreign key (teller_id) references  
colorg.teller not enforced  
DB20000I  The SQL command completed successfully.
```

```
alter table colorg.history add constraint fkeybranch foreign key (branch_id) references  
colorg.branch not enforced  
DB20000I The SQL command completed successfully.
```

```
db2 -tvf runstats.cmd
```

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f rowquery1.sql -I complete -ISO CS | tee rowqlbat.txt  
More rowqlbat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

```
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as  
br_trans  
FROM ROWORG.HISTORY AS HISTORY  
WHERE HISTORY.BRANCH_ID between 10 and 20  
GROUP BY HISTORY.BRANCH_ID  
ORDER BY HISTORY.BRANCH_ID ASC ;
```

BRANCH_ID	BR_BALANCE	BR_TRANS
10	2045913300.00	27777
11	2068728300.00	27544
12	2554543800.00	31732
13	3685920900.00	40134
14	2787550200.00	33372

\* 11 row(s) fetched, 5 row(s) output.

```
* Prepare Time is:      0.205942 seconds  
* Execute Time is:     0.570747 seconds  
* Fetch Time is:       0.005214 seconds  
* Elapsed Time is:     0.781903 seconds (complete)
```



The example elapsed time is 0.781903.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

For example:

```
-----
* Comment: " CHECK STATS"
-----
```

```
* SQL Statement Number 4:
```

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
       ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
 ;
```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
0	1835	297	2132

```
* 1 row(s) fetched, 1 row(s) output.
```

```
* Prepare Time is:      0.106602 seconds
* Execute Time is:      0.031500 seconds
* Fetch Time is:        0.000160 seconds
* Elapsed Time is:      0.138262 seconds (complete)
```

```
-----
* SQL Statement Number 5:
```

```
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
 ;
```

TOTAL_COL_TIME	TOTAL_COMPILE_TIME	POOL_READ_TIME	TOTAL_CPU_TIME
630	344	297	272814

```
* 1 row(s) fetched, 1 row(s) output.
```

```
* Prepare Time is:      0.002589 seconds
* Execute Time is:      0.000295 seconds
* Fetch Time is:        0.000141 seconds
* Elapsed Time is:      0.003025 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (2132 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (650 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (297 in the sample)

Now look at the access plan for the db2exfmt explain report.  
In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o exrowq1.txt
more exrowq1.txt
```

The output should look similar to the following:

Access Plan:

-----

```
Total Cost:      588.169
Query Degree:     1
```

```
      Rows
      RETURN
      (  1)
      Cost
      I/O
      |
      11
      GRPBY
      (  2)
      588.168
      255.345
      |
      11
      TBSCAN
      (  3)
      588.168
      255.345
      |
      11
      SORT
```

```

      (    4)
      588.167
      255.345
      |
      11
      pGRPBY
      (    5)
      588.165
      255.345
      |
      56493.4
      FETCH
      (    6)
      579.791
      255.345
      /---+---\
56493.4      513576
RIDSCN      TABLE: ROWORG
(    7)      HISTORY
116.992      Q1
32.9881
      |
56493.4
SORT
(    8)
116.992
32.9881
      |
56493.4
IXSCAN
(    9)
103.901
32.9881
      |
513576
INDEX: ROWORG
HISTIX1
Q1

```

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (588 in the sample)

Total I/O cost: \_\_\_\_\_ (255 in the sample)

The access plan uses one index to access the table ROWORG.HISTORY.

We will run a similar db2batch report using the same SQL query but the column-organized table CLORORG.HISTORY will be used instead of the row-organized table.

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f colquery1.sql -I complete -ISO CS | tee colqlbat.txt  
more colqlbat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

\* SQL Statement Number 2:

```
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as  
br_trans  
FROM COLORG.HISTORY AS HISTORY  
WHERE HISTORY.BRANCH_ID between 10 and 20  
GROUP BY HISTORY.BRANCH_ID  
ORDER BY HISTORY.BRANCH_ID ASC ;
```

BRANCH_ID	BR_BALANCE	BR_TRANS
10	2045913300.00	27777
11	2068728300.00	27544
12	2554543800.00	31732
13	3685920900.00	40134
14	2787550200.00	33372

\* 11 row(s) fetched, 5 row(s) output.

```
* Prepare Time is:      0.173632 seconds  
* Execute Time is:     0.169731 seconds  
* Fetch Time is:       0.024093 seconds  
* Elapsed Time is:     0.367456 seconds (complete)
```

The example elapsedtime is 0.367456 seconds compared to 0.781903 seconds for the row-organized table based query.

Lookt at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

The output should look similar to the following:

```

-----
* Comment: " CHECK STATS"
-----

* SQL Statement Number 4:

SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
  ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
  from table(mon_get_connection(null,-1)) as con
  where application_name = 'db2batch'
;

POOL_COL_L_READS      POOL_DATA_L_READS      POOL_INDEX_L_READS      TOTAL_L_READS
-----
                        130                        274                        201                        605

* 1 row(s) fetched, 1 row(s) output.

* Prepare Time is:      0.102977 seconds
* Execute Time is:      0.026895 seconds
* Fetch Time is:        0.000187 seconds
* Elapsed Time is:      0.130059 seconds (complete)

-----

* SQL Statement Number 5:

SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
  where application_name = 'db2batch'
;

TOTAL_COL_TIME      TOTAL_COMPILE_TIME      POOL_READ_TIME      TOTAL_CPU_TIME
TOTAL_WAIT_TIME
-----
                        410                        287                        305                        171293
505

* 1 row(s) fetched, 1 row(s) output.

```

```
* Prepare Time is:      0.002879 seconds
* Execute Time is:      0.000368 seconds
* Fetch Time is:        0.000153 seconds
* Elapsed Time is:      0.003400 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (605 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (505 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (305 in the sample)

Now look at the access plan for the db2exfmt explain report.  
In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o excolq1.txt
more excolq1.txt
```

The output should look similar to the following:

Access Plan:

-----

Total Cost:	394.716
Query Degree:	1

```
      Rows
      RETURN
      (  1 )
      Cost
      I/O
      |
      11
      TBSCAN
      (  2 )
      394.716
      44.169
      |
      11
      SORT
      (  3 )
      394.715
      44.169
```

```

      |
      11
    CTQ
    (  4 )
    394.714
    44.169
      |
      11
    GRPBY
    (  5 )
    394.712
    44.169
      |
    56493.4
    TBSCAN
    (  6 )
    393.884
    44.169
      |
    513576
CO-TABLE: COLORG
HISTORY
Q1

```

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (394 in the sample)

Total I/O cost: \_\_\_\_\_ (44 in the sample)

The access plan uses a table scan to access the table COLORG.HISTORY.

**Note:** The estimated I/O cost as well as the actual logical pages read for the query using the column-organized table are less than the query using the row-organized table.

The next SQL query that we will analyze accesses a single table and produces a summarized result from a subset of the table data. In this case the subset of data from the table ROWORG.HISTORY is not based on a column that is indexed. We will use the db2batch application to execute the query and report basic elapsed time and statistics. We ill

also include several SQL queries in the db2batch sequence that return important DB2 performance metrics that we will use to compare the row-organized and column-organized table processing.

We will start by running the SQL query using the row-organized tables. The file *rowquery2.sql* contains the following statements:

```
set current explain mode yes;
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as
br_trans
  FROM ROWORG.HISTORY AS HISTORY
 WHERE HISTORY.acct_ID between 5000 and 80000
 GROUP BY HISTORY.BRANCH_ID
 ORDER BY HISTORY.BRANCH_ID ASC ;

set current explain mode no;
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
  ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as
total_l_reads
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
 ;
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time,
total_wait_time
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
 ;
```

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f rowquery2.sql -I complete -ISO CS | tee rowq2bat.txt
more rowq2bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

\* SQL Statement Number 2:



```
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as
br_trans
FROM ROWORG.HISTORY AS HISTORY
WHERE HISTORY.acct_ID between 5000 and 80000
GROUP BY HISTORY.BRANCH_ID
ORDER BY HISTORY.BRANCH_ID ASC ;
```

BRANCH_ID	BR_BALANCE	BR_TRANS
1	3163900.00	176
2	3964500.00	217
3	3162100.00	174
4	1184000.00	71
5	2442000.00	136

\* 100 row(s) fetched, 5 row(s) output.

```
* Prepare Time is:      0.047030 seconds
* Execute Time is:     0.303496 seconds
* Fetch Time is:       0.004525 seconds
* Elapsed Time is:     0.355051 seconds (complete)
```

The example elapsed time is 0.355051.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

For example:

\* SQL Statement Number 4:

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
from table(mon_get_connection(null,-1)) as con
where application_name = 'db2batch'
;
```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
0	1213	84	1297

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000129 seconds
* Execute Time is:      0.010047 seconds
* Fetch Time is:        0.000222 seconds
* Elapsed Time is:      0.010398 seconds (complete)
```

-----

\* SQL Statement Number 5:

```
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
from table(mon_get_connection(null,-1)) as con
where application_name = 'db2batch'
;
```

TOTAL_COL_TIME	TOTAL_COMPILE_TIME	POOL_READ_TIME	TOTAL_CPU_TIME	TOTAL_WAIT_TIME
270	0	47	0	67514

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000138 seconds
* Execute Time is:      0.000185 seconds
* Fetch Time is:        0.000074 seconds
* Elapsed Time is:      0.000397 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (1297 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (270 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (0 in the sample)

The sample results indicate that the query was able to use data in the buffer pool rather than reading the pages from disk.

Now look at the access plan for the db2exfmt explain report.

In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o exrowq2.txt
more exrowq2.txt
```

The output should look similar to the following:

Access Plan:

-----

Total Cost: 2407.47  
Query Degree: 1

Rows  
RETURN  
( 1)  
Cost  
I/O  
|  
100  
GRPBY  
( 2)  
2407.47  
1118  
|  
100  
TBSCAN  
( 3)  
2407.46  
1118  
|  
100  
SORT  
( 4)  
2407.46  
1118  
|  
100.673  
pGRPBY  
( 5)  
2407.44  
1118  
|  
38521.7  
TBSCAN  
( 6)  
2401.71  
1118  
|  
513576  
TABLE: ROWORG  
HISTORY  
Q1

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (2407 in the sample)

Total I/O cost: \_\_\_\_\_ (1118 in the sample)

The access plan uses a table scan to access the table ROWORG.HISTORY.

We will run a similar db2batch report using the same SQL query but the column-organized table CLORORG.HISTORY will be used instead of the row-organized table.

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f colquery2.sql -I complete -ISO CS | tee colq2bat.txt  
more colq2.bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

\* SQL Statement Number 2:

```
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as  
br_trans  
FROM COLORG.HISTORY AS HISTORY  
WHERE HISTORY.acct_ID between 5000 and 80000  
GROUP BY HISTORY.BRANCH_ID  
ORDER BY HISTORY.BRANCH_ID ASC ;
```

BRANCH_ID	BR_BALANCE	BR_TRANS
1	3163900.00	176
2	3964500.00	217
3	3162100.00	174
4	1184000.00	71
5	2442000.00	136

\* 100 row(s) fetched, 5 row(s) output.

```
* Prepare Time is:      0.012101 seconds
* Execute Time is:      0.114467 seconds
* Fetch Time is:        0.004360 seconds
* Elapsed Time is:      0.130928 seconds (complete)
```

The example elapsedtime is 0.130928 seconds compared to 0.355051 seconds for the row-organized table based query.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

The output should look similar to the following:

```
-----
* Comment: " CHECK STATS"
-----

* SQL Statement Number 4:

SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
       ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
;

POOL_COL_L_READS      POOL_DATA_L_READS      POOL_INDEX_L_READS      TOTAL_L_READS
-----
                148                91                88                327

* 1 row(s) fetched, 1 row(s) output.

* Prepare Time is:      0.000167 seconds
* Execute Time is:      0.009893 seconds
* Fetch Time is:        0.000145 seconds
* Elapsed Time is:      0.010205 seconds (complete)

-----

* SQL Statement Number 5:

SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
;

TOTAL_COL_TIME      TOTAL_COMPILE_TIME      POOL_READ_TIME      TOTAL_CPU_TIME
TOTAL_WAIT_TIME
```

```
-----
-----
192          327          12          35          56207
```

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000133 seconds
* Execute Time is:     0.000171 seconds
* Fetch Time is:       0.000231 seconds
* Elapsed Time is:     0.000535 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (327 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (192 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (35 in the sample)

Now look at the access plan for the db2exfmt explain report.  
In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o excolq2.txt
more excolq2.txt
```

The output should look similar to the following:

Access Plan:

-----

```
Total Cost:      395.575
Query Degree:    1
```

```
    Rows
RETURN
(   1)
    Cost
    I/O
    |
    100
TBSCAN
(   2)
395.575
49.6901
    |
    100
SORT
(   3)
395.572
```

```

49.6901
|
100
CTQ
( 4)
395.555
49.6901
|
100
GRPBY
( 5)
395.549
49.6901
|
38521.7
TBSCAN
( 6)
394.983
49.6901
|
513576
CO-TABLE: COLORG
HISTORY
Q1

```

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (395 in the sample)

Total I/O cost: \_\_\_\_\_ (49 in the sample)

The access plan uses a table scan to access the table COLORG.HISTORY.  
Compare these estimated costs with the explain report based on the row-organized table.

The next SQL query that we will analyze accesses a single table and produces a summarized result from a subset of the table data. In this case the subset of data is based on two predicates that are supported by indexes on the table ROWORG.HISTORY. We will use the db2batch application to execute the query and report basic elapsed time and statistics. We will also include several SQL queries in the db2batch sequence that return important DB2 performance metrics that we will use to compare the row-organized and column-organized table processing.

We will start by running the SQL query using the row-organized tables. The file *rowquery3.sql* contains the following statements:

```
set current explain mode yes;
```

```
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as  
br_trans  
FROM ROWORG.HISTORY AS HISTORY  
where branch_id between 10 and 20 and teller_id > 800  
GROUP BY HISTORY.BRANCH_ID  
ORDER BY HISTORY.BRANCH_ID ASC ;
```

```
set current explain mode no;  
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,  
  ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as  
total_l_reads  
  from table(mon_get_connection(null,-1)) as con  
  where application_name = 'db2batch'  
  ;  
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time,  
total_wait_time  
  from table(mon_get_connection(null,-1)) as con  
  where application_name = 'db2batch'  
  ;
```

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f rowquery3.sql -I complete -ISO CS | tee rowq3bat.txt  
more rowq3bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

\* SQL Statement Number 2:

```
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as  
br_trans  
FROM ROWORG.HISTORY AS HISTORY  
where branch_id between 10 and 20 and teller_id > 800  
GROUP BY HISTORY.BRANCH_ID  
ORDER BY HISTORY.BRANCH_ID ASC ;
```



BRANCH_ID	BR_BALANCE	BR_TRANS
10	273951900.00	4521
11	684840000.00	7459
12	222541400.00	4126
13	478056000.00	6117
14	220693100.00	4266

\* 11 row(s) fetched, 5 row(s) output.

\* Prepare Time is: 0.048327 seconds  
 \* Execute Time is: 0.547737 seconds  
 \* Fetch Time is: 0.007361 seconds  
 \* Elapsed Time is: 0.603425 seconds (complete)

The example elapsed time is 0.603425.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

For example:

\* SQL Statement Number 4:

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
      ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
from table(mon_get_connection(null,-1)) as con
where application_name = 'db2batch'
;
```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
0	1671	272	1943

\* 1 row(s) fetched, 1 row(s) output.

\* Prepare Time is: 0.020876 seconds  
 \* Execute Time is: 0.010026 seconds  
 \* Fetch Time is: 0.000147 seconds  
 \* Elapsed Time is: 0.031049 seconds (complete)

\* SQL Statement Number 5:

```
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
  where application_name = 'db2batch'
;
```

TOTAL_COL_TIME	TOTAL_COMPILE_TIME	POOL_READ_TIME	TOTAL_CPU_TIME	TOTAL_WAIT_TIME
426	0	81	52	166897

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.002257 seconds
* Execute Time is:      0.000268 seconds
* Fetch Time is:        0.000124 seconds
* Elapsed Time is:      0.002649 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (1943 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (426 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (52 in the sample)

The sample results indicate that the query was able to use data in the buffer pool rather than reading the pages from disk.

Now look at the access plan for the db2exfmt explain report.

In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o exrowq3.txt
more exrowq3.txt
```

The output should look similar to the following:

```
Access Plan:
-----
      Total Cost:      551.41
      Query Degree:    1

      Rows
```

```

RETURN
( 1)
Cost
I/O
|
11
GRPBY
( 2)
551.41
227.047
|
11
TBSCAN
( 3)
551.409
227.047
|
11
SORT
( 4)
551.409
227.047
|
11
pGRPBY
( 5)
551.407
227.047
|
11298.7
FETCH
( 6)
549.718
227.047
/---+---\
11298.7      513576
RIDSCN      TABLE: ROWORG
( 7)        HISTORY
217.377      Q1
63.2933
|
11298.7
SORT
( 8)
217.377
63.2933
|
11298.7

```

IXAND	
( 9)	
214.328	
63.2933	
/-----+-----\	
56493.4	102715
IXSCAN	IXSCAN
( 10)	( 11)
103.901	106.627
32.9881	30.3053
513576	513576
INDEX: ROWORG	INDEX: ROWORG
HISTIX1	HISTIX2
Q1	Q1

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (551 in the sample)

Total I/O cost: \_\_\_\_\_ (227 in the sample)

The access plan uses the index anding IXAND operation to combine two index scan results to access the table ROWORG.HISTORY.

We will run a similar db2batch report using the same SQL query but the column-organized table CLORORG.HISTORY will be used instead of the row-organized table.

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f colquery2.sql -I complete -ISO CS | tee colq3bat.txt
more colq3bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

\* SQL Statement Number 2:

```
SELECT HISTORY.BRANCH_ID, sum(HISTORY.balance) as br_balance, count(*) as
br_trans
FROM COLORG.HISTORY AS HISTORY
```

```
where branch_id between 10 and 20 and teller_id > 800
GROUP BY HISTORY.BRANCH_ID
ORDER BY HISTORY.BRANCH_ID ASC ;
```

BRANCH_ID	BR_BALANCE	BR_TRANS
10	273951900.00	4521
11	684840000.00	7459
12	222541400.00	4126
13	478056000.00	6117
14	220693100.00	4266

\* 11 row(s) fetched, 5 row(s) output.

```
* Prepare Time is:      0.040934 seconds
* Execute Time is:     0.049338 seconds
* Fetch Time is:       0.023248 seconds
* Elapsed Time is:     0.113520 seconds (complete)
```

The example elapsedtime is 0.113520 seconds compared to 0.603425 seconds for the row-organized table based query.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

The output should look similar to the following:

\* SQL Statement Number 4:

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
from table(mon_get_connection(null,-1)) as con
where application_name = 'db2batch'
;
```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
157	98	90	345

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000121 seconds
* Execute Time is:     0.010118 seconds
* Fetch Time is:       0.000180 seconds
* Elapsed Time is:     0.010419 seconds (complete)
```

-----  
\* SQL Statement Number 5:

```
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
  where application_name = 'db2batch'
;
```

TOTAL_COL_TIME	TOTAL_COMPILE_TIME	POOL_READ_TIME	TOTAL_CPU_TIME	TOTAL_WAIT_TIME
113	138	41	6	42409

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000136 seconds
* Execute Time is:      0.000177 seconds
* Fetch Time is:        0.000071 seconds
* Elapsed Time is:      0.000384 seconds (complete)
```

Note the following statistics from your copy of the report:

```
TOTAL_L_READS:_____ (90 in the sample)
TOTAL_WAIT_TIME:_____ (113 in the sample)
POOL_READ_TIME:_____ (6 in the sample)
```

Now look at the access plan for the db2exfmt explain report.  
In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o excolq3.txt
more excolq3.txt
```

The output should look similar to the following:

```
Access Plan:
-----
      Total Cost:      400.431
      Query Degree:    1

      Rows
      RETURN
      ( 1)
      Cost
```

```

      I/O
      |
      11
    TBSCAN
    (   2)
    400.431
    49.6901
      |
      11
    SORT
    (   3)
    400.431
    49.6901
      |
      11
    CTQ
    (   4)
    400.429
    49.6901
      |
      11
    GRPBY
    (   5)
    400.428
    49.6901
      |
    11299
    TBSCAN
    (   6)
    400.262
    49.6901
      |
    513576
CO-TABLE: COLORG
HISTORY
Q1

```

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (400 in the sample)

Total I/O cost: \_\_\_\_\_ (49 in the sample)

The access plan uses a table scan to access the table COLORG.HISTORY.  
Compare these estimated costs with the explain report based on the row-organized table.

**Note:** We have used three SQL queries to compare performance of row-organized and column-organized tables. For the row-organized table, one query did not use an index, one used a single index and one used multiple indexes. In the sample results, the column-organized table was able to produce the same results faster with no traditional indexes.

The next SQL query that we will analyze joins two tables to produce a summarized result from a subset of the table data. In this case the subset of data is based on one predicate on the larger table that is supported by an index on the table ROWORG.HISTORY. We will use the db2batch application to execute the query and report basic elapsed time and statistics. We will also include several SQL queries in the db2batch sequence that return important DB2 performance metrics that we will use to compare the row-organized and column-organized table processing.

We will start by running the SQL query using the row-organized tables. The file *rowquery4.sql* contains the following statements:

```
set current explain mode yes;
```

```
SELECT Teller.TELLER_ID, sum(HISTORY.BALANCE) as total_balance,  
Teller.TELLER_NAME , count(*) as transactions  
FROM ROWORG.HISTORY AS HISTORY, ROWORG.TELLER AS TELLER  
WHERE HISTORY.TELLER_ID = TELLER.TELLER_ID  
and teller.teller_id between 10 and 30  
GROUP BY TELLER.TELLER_ID, TELLER.TELLER_NAME  
order by 4 desc ;
```

```
set current explain mode no;
```

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,  
( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as  
total_l_reads
```

```
from table(mon_get_connection(null,-1)) as con  
where application_name = 'db2batch'
```

```
;
```

```
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time,  
total_wait_time
```

```
from table(mon_get_connection(null,-1)) as con  
where application_name = 'db2batch'
```

```
;
```

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f rowquery4.sql -I complete -ISO CS | tee rowq4bat.txt
```



```
more rowq4bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

\* SQL Statement Number 2:

```
SELECT Teller.TELLER_ID, sum(HISTORY.BALANCE) as total_balance,
TELLER.TELLER_NAME , count(*) as transactions
FROM ROWORG.HISTORY AS HISTORY, ROWORG.TELLER AS TELLER
WHERE HISTORY.TELLER_ID = TELLER.TELLER_ID
and teller.teller_id between 10 and 30
GROUP BY TELLER.TELLER_ID, TELLER.TELLER_NAME
order by 4 desc ;
```

TELLER_ID	TOTAL_BALANCE	TELLER_NAME
2078	273968000.00	2078
1706	197057800.00	1706
426	6940400.00	426
403	6447500.00	403
385	5705900.00	385

\* 21 row(s) fetched, 5 row(s) output.

```
* Prepare Time is:      0.095087 seconds
* Execute Time is:     0.176888 seconds
* Fetch Time is:       0.002835 seconds
* Elapsed Time is:     0.274810 seconds (complete)
```

The example elapsed time is 0.278410.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

For example:

\* SQL Statement Number 4:

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
      ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
;
```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
0	1690	185	1875

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000125 seconds
* Execute Time is:      0.010717 seconds
* Fetch Time is:        0.000150 seconds
* Elapsed Time is:      0.010992 seconds (complete)
```

\* SQL Statement Number 5:

```
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
;
```

TOTAL_COL_TIME	TOTAL_COMPILE_TIME	POOL_READ_TIME	TOTAL_CPU_TIME
245	95	51	38732

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000216 seconds
* Execute Time is:      0.000187 seconds
```

```
* Fetch Time is:          0.000085 seconds
* Elapsed Time is:       0.000488 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (1875 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (245 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (51 in the sample)

The sample results indicate that the query was able to use data in the buffer pool rather than reading the pages from disk.

Now look at the access plan for the db2exfmt explain report.

In the terminal session issue the commands:

```
db2exfmt -l -d db2blu -o exrowq4.txt
more exrowq4.txt.
```

The output should look similar to the following:

```
Access Plan:
-----
      Total Cost:          1558.41
      Query Degree:       1

      Rows
      RETURN
      ( 1)
      Cost
      I/O
      |
      21
      TBSCAN
      ( 2)
      1558.41
      760.626
      |
      21
      SORT
      ( 3)
      1558.41
      760.626
      |
      21
      GRPBY
```

```

      (  4)
    1558.41
    760.626
      |
      21
    TBSCAN
    (   5)
    1558.4
    760.626
      |
      21
    SORT
    (   6)
    1558.4
    760.626
      |
    10785.1
    ^HSJOIN
    (   7)
    1556.48
    760.626
      /-----+-----\
    10785.1                                21
    FETCH                                FETCH
    (   8)                                (  12)
    1549.15                                7.08178
    759.626                                1
      /---+---\                          /---+---\
    10785.1      513576                  21      1000
    RIDSCAN      TABLE: ROWORG          IXSCAN      TABLE: ROWORG
    (   9)        HISTORY                  (  13)        TELLER
    25.7899          Q2                   0.0327543      Q1
    3.045
      |
    10785.1
    SORT
    (  10)
    25.7898
    3.045
      |
    10785.1
    IXSCAN
    (  11)
    23.7304
    3.045
      |
    513576
    INDEX: ROWORG
    INDEX: SYSIBM
    SQL131107105256150
      Q1

```

HISTIX2  
Q2

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (1558 in the sample)

Total I/O cost: \_\_\_\_\_ (760 in the sample)

The access plan uses the indexes on both tables. The join processing uses a hash join method.

We will run a similar db2batch report using the same SQL query but the column-organized table CLORORG.HISTORY will be used instead of the row-organized table.

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f colquery4.sql -I complete -ISO CS | tee colq4bat.txt
more colq4bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

-----

\* SQL Statement Number 2:

```
SELECT Teller.TELLER_ID, sum(HISTORY.BALANCE) as total_balance,
TELLER.TELLER_NAME , count(*) as transactions
FROM COLORG.HISTORY AS HISTORY, COLORG.TELLER AS TELLER
WHERE HISTORY.TELLER_ID = TELLER.TELLER_ID
and teller.teller_id between 10 and 30
GROUP BY TELLER.TELLER_ID, TELLER.TELLER_NAME
order by 4 desc ;
```

```
TELLER_ID TOTAL_BALANCE          TELLER_NAME
TRANSACTIONS
```

-----  
-

```

      30                273968000.00 <--20 BYTE STRING-->
2078
      16                197057800.00 <--20 BYTE STRING-->
1706
      25                6940400.00 <--20 BYTE STRING-->
426
      28                6447500.00 <--20 BYTE STRING-->
403
      10                5705900.00 <--20 BYTE STRING-->
385

```

\* 21 row(s) fetched, 5 row(s) output.

```

* Prepare Time is:      0.012669 seconds
* Execute Time is:     0.081109 seconds
* Fetch Time is:       0.016374 seconds
* Elapsed Time is:     0.110152 seconds (complete)

```

The example elapsedtime is 0.110152 seconds compared to 0.274810 seconds for the row-organized table based query.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

The output should look similar to the following:

\* SQL Statement Number 4:

```

SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
      ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
;

```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
-----	-----	-----	-----
160	165	144	469

\* 1 row(s) fetched, 1 row(s) output.

```

* Prepare Time is:      0.018891 seconds
* Execute Time is:     0.045660 seconds
* Fetch Time is:       0.000256 seconds
* Elapsed Time is:     0.064807 seconds (complete)

```

```

-----

* SQL Statement Number 5:

SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
  where application_name = 'db2batch'
;

TOTAL_COL_TIME          TOTAL_COMPILE_TIME    POOL_READ_TIME          TOTAL_CPU_TIME
TOTAL_WAIT_TIME
-----
156                    238                    32                    47                    87176

* 1 row(s) fetched, 1 row(s) output.

* Prepare Time is:      0.000422 seconds
* Execute Time is:      0.000454 seconds
* Fetch Time is:        0.000187 seconds
* Elapsed Time is:      0.001063 seconds (complete)

```

Note the following statistics from your copy of the report:

```

TOTAL_L_READS:_____ (469 in the sample)
TOTAL_WAIT_TIME:_____ (238 in the sample)
POOL_READ_TIME:_____ (47 in the sample)

```

In the sample report there are nearly as many index object pages with logical reads as column-organized object and data object page reads.

Column-organized tables do have the page map indexes which are used internally to locate the pages allocated for each column, and these references are counted as index page reads. The page map indexes do not appear in the access plans in explain reports.

Now look at the access plan for the db2exfmt explain report.

In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o excolq4.txt
more excolq4.txt
```

The output should look similar to the following:

```

Access Plan:
-----
Total Cost:      416.306

```

Query Degree: 1

```

      Rows
      RETURN
      ( 1 )
      Cost
      I/O
      |
      21.0153
      TBSCAN
      ( 2 )
      416.306
      46.6953
      |
      21.0153
      SORT
      ( 3 )
      416.306
      46.6953
      |
      21.0153
      CTQ
      ( 4 )
      416.303
      46.6953
      |
      21.0153
      ^HSJOIN
      ( 5 )
      416.302
      46.6953
      /-----+-----\
      1000          20.679
      TBSCAN      GRPBY
      ( 6 )      ( 7 )
      17.8417    398.455
      2.52632    44.169
      |          |
      1000      10792.9
      CO-TABLE: COLORG  TBSCAN
      TELLER      ( 8 )
      Q1          398.296
              44.169
              |
              513576
      CO-TABLE: COLORG
      HISTORY
      Q2
  
```



Review the db2exfmt report noting the following information:

Total estimated cost:\_\_\_\_\_ (416 in the sample)

Total I/O cost:\_\_\_\_\_ (46 in the sample)

The access plan uses a table scan to access the tables COLORG.HISTORY and COLORG.TELLER.  
Compare these estimated costs with the explain report based on the row-organized table.

Locate the detail for GROUP BY (GRPBY) operation number 7 in the explain report. It will look similar to the following.

For example:

```

7) GRPBY : (Group By)
    Cumulative Total Cost:          398.455
    Cumulative CPU Cost:           4.49065e+08
    Cumulative I/O Cost:           44.169
    Cumulative Re-Total Cost:    54.78
    Cumulative Re-CPU Cost:      4.48934e+08
    Cumulative Re-I/O Cost:       0
    Cumulative First Row Cost:  7.25177
    Estimated Bufferpool Buffers:    0

Arguments:
-----
AGGMODE : (Aggregation Mode)
        HASHED COMPLETE
GROUPBYC: (Group By columns)
        TRUE
GROUPBYN: (Number of Group By columns)
        1
GROUPBYR: (Group By requirement)
        1: Q3.TELLER_ID
JN INPUT: (Join input leg)
        INNER

Input Streams:
-----
    4) From Operator #8

        Estimated number of rows:  10792.9
        Number of columns:         2
        Subquery predicate ID:           Not Applicable

        Column Names:
        -----
        +Q3.TELLER_ID+Q3.BALANCE

```

Output Streams:

-----

5) To Operator #5

Estimated number of rows: 20.679  
Number of columns: 3  
Subquery predicate ID: Not Applicable

Column Names:

-----

+Q4.\$C2+Q4.\$C1+Q4.TELLER\_ID

The aggregation mode is 'Hashed Complete', which uses the sort heap memory to perform the grouping without the need to sort the input first. The small summarized result then becomes the inner table for the hash join operation.

The next SQL query that we will analyze joins three tables to produce a composite result from a subset of the table data. The query includes predicates on two of the tables, ROWORG.HISTORY and ROWORG.ACCT. the tables are joined using columns that are unique indexes of the ACCT and TELLER tables. The SQL query contains an ORDER BY clause, but no GROUP BY clause. We will use the db2batch application to execute the query and report basic elapsed time and statistics. We will also include several SQL queries in the db2batch sequence that return important DB2 performance metrics that we will use to compare the row-organized and column-organized table processing.

We will start by running the SQL query using the row-organized tables. The file *rowquery5.sql* contains the following statements:

```
set current explain mode yes;
```

```
SELECT ACCT.ACCT_ID, ACCT.NAME, TELLER.TELLER_ID, TELLER.TELLER_NAME,  
       TELLER.TELLER_CODE, ACCT.ADDRESS ,  
       HISTORY.BRANCH_ID, HISTORY.BALANCE, HISTORY.PID, HISTORY.TEMP  
FROM roworg.ACCT AS ACCT, roworg.TELLER AS TELLER, roworg.HISTORY AS  
HISTORY  
WHERE ACCT.ACCT_ID = HISTORY.ACCT_ID  
AND ACCT.ACCT_GRP BETWEEN 100 AND 200  
AND HISTORY.TELLER_ID = TELLER.TELLER_ID  
AND HISTORY.BRANCH_ID BETWEEN 40 AND 50  
ORDER BY HISTORY.PID ASC ;
```

```
set current explain mode no;
```

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
      ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as
total_l_reads
      from table(mon_get_connection(null,-1)) as con
      where application_name = 'db2batch'
;

SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time,
total_wait_time
      from table(mon_get_connection(null,-1)) as con
      where application_name = 'db2batch'
;
```

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f rowquery5.sql -I complete -ISO CS | tee rowq5bat.txt
more rowq5bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

\* SQL Statement Number 2:

```
SELECT ACCT.ACCT_ID, ACCT.NAME, TELLER.TELLER_ID, TELLER.TELLER_NAME,
      TELLER.TELLER_CODE, ACCT.ADDRESS ,
      HISTORY.BRANCH_ID, HISTORY.BALANCE, HISTORY.PID, HISTORY.TEMP
      FROM roworg.ACCT AS ACCT, roworg.TELLER AS TELLER, roworg.HISTORY AS HISTORY
      WHERE ACCT.ACCT_ID = HISTORY.ACCT_ID
      AND ACCT.ACCT_GRP BETWEEN 100 AND 200
      AND HISTORY.TELLER_ID = TELLER.TELLER_ID
      AND HISTORY.BRANCH_ID BETWEEN 40 AND 50
      ORDER BY HISTORY.PID ASC ;
```

ACCT_ID	NAME	TELLER_ID	TELLER_NAME	TELLER_CODE	ADDRESS
BRANCH_ID	BALANCE	PID	TEMP		
9517	<--20 BYTE STRING-->	893	<--20 BYTE STRING-->	ab	<----- 30 BYTE
STRING	----->	45	26100.00	0 TP1ST	
9517	<--20 BYTE STRING-->	893	<--20 BYTE STRING-->	ab	<----- 30 BYTE
STRING	----->	45	17700.00	0 TP1ST	

```

15393 <--20 BYTE STRING-->      263 <--20 BYTE STRING--> ab      <----- 30 BYTE
STRING ----->      40      27600.00      0 TP1ST
15393 <--20 BYTE STRING-->      263 <--20 BYTE STRING--> ab      <----- 30 BYTE
STRING ----->      40      16700.00      0 TP1ST
15393 <--20 BYTE STRING-->      263 <--20 BYTE STRING--> ab      <----- 30 BYTE
STRING ----->      40      24800.00      0 TP1ST

```

\* 1874 row(s) fetched, 5 row(s) output.

```

* Prepare Time is:      0.063917 seconds
* Execute Time is:      0.334208 seconds
* Fetch Time is:        0.005953 seconds
* Elapsed Time is:      0.404078 seconds (complete)

```

-----

The example elapsed time is 0.404078.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.

For example:

\* SQL Statement Number 4:

```

SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
  ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
  from table(mon_get_connection(null,-1)) as con
  where application_name = 'db2batch'
;

```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
0	1364	193	1557

\* 1 row(s) fetched, 1 row(s) output.

```

* Prepare Time is:      0.000185 seconds
* Execute Time is:      0.010261 seconds
* Fetch Time is:        0.000148 seconds
* Elapsed Time is:      0.010594 seconds (complete)

```

-----

\* SQL Statement Number 5:

```

SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time

```

```
from table(mon_get_connection(null,-1)) as con
  where application_name = 'db2batch'
;
```

TOTAL_COL_TIME	TOTAL_COMPILE_TIME	POOL_READ_TIME	TOTAL_CPU_TIME
TOTAL_WAIT_TIME			
-----	-----	-----	-----
-----	0	64	53
152			224943

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000128 seconds
* Execute Time is:      0.000170 seconds
* Fetch Time is:        0.000085 seconds
* Elapsed Time is:      0.000383 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (1557 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (152 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (53 in the sample)

The sample results indicates that the query was able to use data in the buffer pool rather than reading the pages from disk.

Now look at the access plan for the db2exfmt explain report.

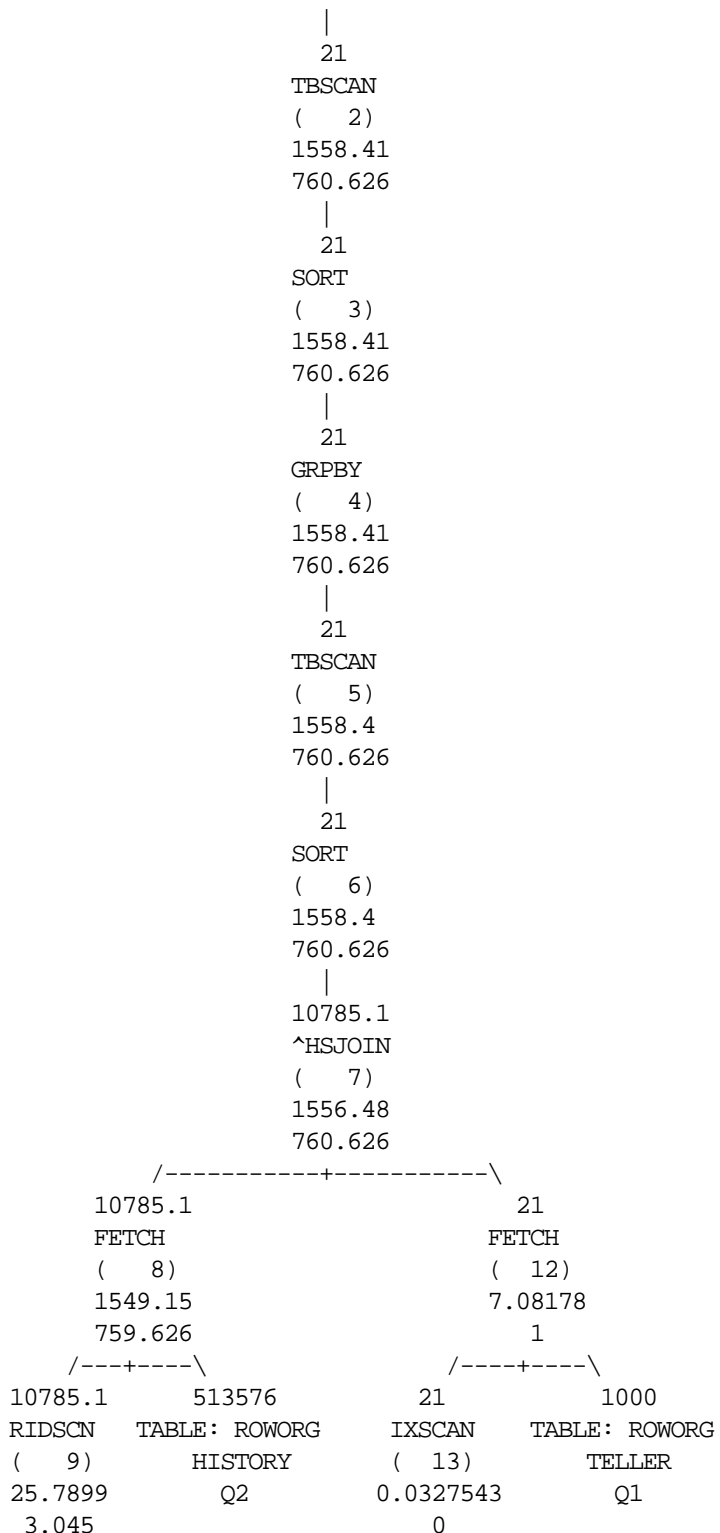
In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o exrowq5.txt
more exrowq5.txt
```

The output should look similar to the following:

```
Access Plan:
-----
Total Cost:      1558.41
Query Degree:    1

              Rows
              RETURN
              (  1 )
              Cost
              I/O
```



10785.1	1000
SORT	INDEX: SYSIBM
( 10)	SQL131107105256150
25.7898	Q1
3.045	
10785.1	
IXSCAN	
( 11)	
23.7304	
3.045	
513576	
INDEX: ROWORG	
HISTIX2	
Q2	

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (2135 in the sample)

Total I/O cost: \_\_\_\_\_ (832 in the sample)

The access plan uses the indexes on the HISTORY and TELLER tables, but performs a table scan on the ACCT table, since the predicate on the ACCT\_GRP column is not supported by an index. The join processing utilizes two hash join methods.

We will run a similar db2batch report using the same SQL query but the column-organized table COLORG.HISTORY will be used instead of the row-organized table.

In the Linux terminal session, issue the command:

```
db2batch -d db2blu -f colquery5.sql -I complete -ISO CS | tee colq5bat.txt
more colq5bat.txt
```

Review the db2batch report noting the following information

Elapsed time for the query (Statement2): \_\_\_\_\_

For example:

-----

\* SQL Statement Number 2:

```
SELECT ACCT.ACCT_ID, ACCT.NAME, TELLER.TELLER_ID, TELLER.TELLER_NAME,
       TELLER.TELLER_CODE, ACCT.ADDRESS ,
       HISTORY.BRANCH_ID, HISTORY.BALANCE, HISTORY.PID, HISTORY.TEMP
FROM colorg.ACCT AS ACCT, colorg.TELLER AS TELLER, colorg.HISTORY AS
HISTORY
WHERE ACCT.ACCT_ID = HISTORY.ACCT_ID
AND ACCT.ACCT_GRP BETWEEN 100 AND 200
AND HISTORY.TELLER_ID = TELLER.TELLER_ID
AND HISTORY.BRANCH_ID BETWEEN 40 AND 50
ORDER BY HISTORY.PID ASC ;
```

ACCT_ID	NAME	TELLER_ID	TELLER_NAME	TELLER_CODE
ADDRESS		BRANCH_ID	BALANCE	PID
305649	<--20 BYTE STRING-->	147	<--20 BYTE STRING-->	ab
<----- 30 BYTE STRING ----->		41	10100.00	0 TP1ST
144705	<--20 BYTE STRING-->	899	<--20 BYTE STRING-->	ab
<----- 30 BYTE STRING ----->		48	10000.00	0 TP1ST
609072	<--20 BYTE STRING-->	398	<--20 BYTE STRING-->	ab
<----- 30 BYTE STRING ----->		50	10000.00	0 TP1ST
122255	<--20 BYTE STRING-->	623	<--20 BYTE STRING-->	ab
<----- 30 BYTE STRING ----->		43	10100.00	0 TP1ST
184091	<--20 BYTE STRING-->	972	<--20 BYTE STRING-->	ab
<----- 30 BYTE STRING ----->		42	10000.00	0 TP1ST

\* 1874 row(s) fetched, 5 row(s) output.

```
* Prepare Time is:      0.013802 seconds
* Execute Time is:     0.376213 seconds
* Fetch Time is:       0.008739 seconds
* Elapsed Time is:     0.398754 seconds (complete)
```

The example elapsedtime is 0.398754 seconds compared to 0.404078 seconds for the row-organized table based query.

Look at the two queries that return performance metrics using the table function MON\_GET\_CONNECTION.



The output should look similar to the following:

\* SQL Statement Number 4:

```
SELECT pool_col_l_reads, pool_data_l_reads, pool_index_l_reads ,
      ( pool_col_l_reads + pool_data_l_reads + pool_index_l_reads ) as total_l_reads
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
;
```

POOL_COL_L_READS	POOL_DATA_L_READS	POOL_INDEX_L_READS	TOTAL_L_READS
538	178	169	885

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000126 seconds
* Execute Time is:      0.010261 seconds
* Fetch Time is:        0.000146 seconds
* Elapsed Time is:      0.010533 seconds (complete)
```

\* SQL Statement Number 5:

```
SELECT total_col_time, total_compile_time , pool_read_time, total_cpu_time, total_wait_time
  from table(mon_get_connection(null,-1)) as con
 where application_name = 'db2batch'
;
```

TOTAL_COL_TIME	TOTAL_COMPILE_TIME	POOL_READ_TIME	TOTAL_CPU_TIME	TOTAL_WAIT_TIME
939	14	77	227041	472

\* 1 row(s) fetched, 1 row(s) output.

```
* Prepare Time is:      0.000133 seconds
* Execute Time is:      0.000169 seconds
* Fetch Time is:        0.000080 seconds
* Elapsed Time is:      0.000382 seconds (complete)
```

Note the following statistics from your copy of the report:

TOTAL\_L\_READS:\_\_\_\_\_ (885 in the sample)

TOTAL\_WAIT\_TIME:\_\_\_\_\_ (472 in the sample)

POOL\_READ\_TIME:\_\_\_\_\_ (77 in the sample)

In the sample report the increase in column-oriented object page references relates to the access to the ACCT table, which is relatively large. This increases the pool read time and total wait time.

Now look at the access plan for the db2exfmt explain report.

In the terminal session issue the commands:

```
db2exfmt -1 -d db2blu -o excolq5.txt
more excolq5.txt
```

The output should look similar to the following:

Access Plan:

-----

Total Cost:	784.282
Query Degree:	1

Rows
RETURN
( 1 )
Cost
I/O
5863.4
TBSCAN
( 2 )
784.282
152.181
5863.4
SORT
( 3 )
784.103
152.181
5863.4
CTQ
( 4 )
782.339
152.181
5863.4
^HSJOIN
( 5 )
782.261
152.181

```

          /-----+-----\
        5769.59          1000
        ^HSJOIN          TBSCAN
        (  6)          (  9)
        763.274          18.9541
        149.497          2.68421
          /-----+-----\
        56493.4          102129          1000
        TBSCAN          TBSCAN    CO-TABLE: COLORG
        (  7)          (  8)          TELLER
        412.349          350.06          Q2
        96.6197          52.8772
        |                |
        513576          1e+06
    CO-TABLE: COLORG    CO-TABLE: COLORG
        HISTORY          ACCT
        Q1                Q3

```

Review the db2exfmt report noting the following information:

Total estimated cost: \_\_\_\_\_ (784 in the sample)

Total I/O cost: \_\_\_\_\_ (152 in the sample)

These estimated costs are less than the costs estimated to access the row-organized tables.

**Note:** Your results may have been different, but in each case the column-organized tables used the DB2 BLU Acceleration support to return results with good performance and reduced system resources compared to a set of standard row-organized tables.

We could have performed some additional tuning with the row-organized tables, looking for better indexes or table organizations, but the column-organized tables may offer a simple solution that performs well and does not require the additional performance analysis.

Use the command file named db2bluend to reset the DB2 Registry variable DB2\_WORKLOAD and to drop the test database DB2BLU.

In the terminal session, issue the command:

```
./db2bluend
```

Close all remaining terminal sessions now.

For each remaining terminal sessions:

```
db2 terminate
exit
```

This concludes the second part of BLU Acceleration exercise.



---

© Copyright IBM Corporation 2014  
All Rights Reserved.

IBM Canada  
8200 Warden Avenue  
Markham, ON  
L6G 1C7  
Canada

IBM, the IBM logo, ibm.com and Tivoli are trademarks or registered trademarks of International Business Machines Corporation in the United States, other countries, or both. If these and other IBM trademarked terms are marked on their first occurrence in this information with a trademark symbol (® or ™), these symbols indicate U.S. registered or common law trademarks owned by IBM at the time this information was published. Such trademarks may also be registered or common law trademarks in other countries. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at [ibm.com/legal/copytrade.shtml](http://ibm.com/legal/copytrade.shtml)

Other company, product and service names may be trademarks or service marks of others.

References in this publication to IBM products and services do not imply that IBM intends to make them available in all countries in which IBM operates.

No part of this document may be reproduced or transmitted in any form without written permission from IBM Corporation.

Product data has been reviewed for accuracy as of the date of initial publication. Product data is subject to change without notice. Any statements regarding IBM's future direction and intent are subject to change or withdrawal without notice, and represent goals and objectives only.

THE INFORMATION PROVIDED IN THIS DOCUMENT IS DISTRIBUTED "AS IS" WITHOUT ANY WARRANTY, EITHER EXPRESS OR IMPLIED. IBM EXPRESSLY DISCLAIMS ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE OR NON-INFRINGEMENT.

IBM products are warranted according to the terms and conditions of the agreements (e.g. IBM Customer Agreement, Statement of Limited Warranty, International Program License Agreement, etc.) under which they are provided.